

Beginner Homelab on a Raspberry Pi

Self-host your media, block ads network-wide, get a one-URL control panel, reach it from mobile devices, and route traffic through a VPN, all over a private Tailscale network

Evanns Morales-Cuadrado

2026-06-02

Table of contents

Preface. What this series builds	5
Volume I: Building the homelab	5
Volume II: VPNs and remote access	5
Volume III: Extras	6
The hardware and accounts	6
Conventions used throughout	6
I Volume I: Building the homelab	8
1 Chapter 1. Foundation: your Raspberry Pi and your Tailscale network	9
1.1 Why Tailscale is the spine of everything	9
1.2 Prerequisites	9
1.3 Part A: Flash and first boot	10
1.4 Part B: Join the tailnet	11
1.5 Part C: Add your devices and verify	12
1.6 Recap	13
2 Chapter 2. Streaming audiobooks to a phone with Audiobookshelf	14
2.1 Prerequisites	14
2.2 Part A: Rip and copy the audio	15
2.3 Part B: Run the server	17
2.4 Part C: Connect and use it	18
2.5 Maintenance	20
2.6 Troubleshooting	20
2.7 Recap	21
3 Chapter 3. Network-wide ad blocking at home with Pi-hole	22
3.1 What you are building (and why this design)	22
3.2 Part A: Run Pi-hole on the Pi	22
3.3 Part B: Point your home network at it	27
3.4 Part C: Add blocklists and verify	28
3.5 Troubleshooting	29
3.6 Recap	30
4 Chapter 4. Pretty URLs: a reverse proxy + local DNS	31
4.1 How this works: two jobs, two tools	31
4.2 Part A: Wire the services onto one network	32
4.3 Part B: Deploy Caddy and DNS	34
4.4 Part C: Use it and trust the cert	37
4.5 The pattern you'll reuse	38

4.6	Caveats	39
4.7	Recap	39
5	Chapter 5. A one-URL dashboard at <code>https://home.home</code>	40
5.1	Part A: Deploy Homepage and Portainer	40
5.2	Part B: Give it a pretty URL and use it	45
5.3	Troubleshooting	46
5.4	Recap	46
II	Volume II: VPNs and remote access	48
6	Chapter 6. Why a VPN, and Mullvad vs. a standard app like Aura	49
6.1	What a VPN does (and what it doesn't)	49
6.2	The standard consumer VPN (Aura, and apps like it)	49
6.3	Mullvad: the privacy-first alternative	50
6.4	Which one should you want?	51
6.5	Recap	51
7	Chapter 7. Exit nodes, the Mullvad add-on, and what I run	52
7.1	Reaching your homelab from anywhere (the easy part)	52
7.2	Pi-hole already follows you (set up in Chapter 4)	52
7.3	The one rule that governs everything off-network	53
7.4	What I run: keep them separate, switch per activity	53
7.5	The thing that <i>would</i> let you have both: an exit node	54
7.6	Flavor 1 (the upgrade path): the Tailscale Mullvad add-on	55
7.7	Flavor 2: your own Pi as an exit node (free, but not private)	56
7.8	Flavor 3 (advanced): a Pi-style exit node chained through Mullvad	56
7.9	A router-level VPN might be another answer, but I can't test it yet	57
7.10	The decision matrix	57
7.11	A practical everyday routine	58
7.12	Recap: the whole series	58
III	Volume III: Extras	60
8	Chapter 8. Remoting into machines from a phone	61
8.1	Part A. SSH from a phone with Termius	61
8.2	Part B. A full desktop with NoMachine	62
8.3	Recap	63
9	Chapter 9. Wiring a phone into Linux machines	64
9.1	The two Wi-Fi tools (and why I start with LocalSend)	64
9.2	Install them	65
9.3	Making it work <i>anywhere</i> : over Tailscale	67
9.4	Phone the headless Pi	67
9.5	The wired fallback: USB access with libimobiledevice	68
9.6	Where to start	68
9.7	Recap: and the series, complete	69

Appendices	70
A Appendix A. The Docker Compose files, line by line	70
A.1 Compose concepts that show up everywhere	70
A.2 ~/pihole/compose.yaml: Pi-hole	71
A.3 ~/caddy/compose.yaml + Caddyfile: the reverse proxy	73
A.4 ~/dashboard/compose.yaml: Homepage + Portainer	74
A.5 ~/audiobookshelf/compose.yaml: Audiobookshelf	76
B Appendix B. What should just work, and how to verify it	78
B.1 The one-time cloud setting everything depends on	78
B.2 The service map	78
B.3 Verify it, layer by layer	79
B.4 What works away from home (and what deliberately doesn't)	80
B.5 Reboot behavior	80

Preface. What this series builds

This is my public homelab guide: one Raspberry Pi, one private [Tailscale](#) network, and a handful of useful services added one at a time. The specifics come from the hardware and devices I have access to, but the pattern is meant for anyone building a similar beginner homelab. No port forwarding, no public dashboard hanging off the internet, no mystery pile of commands you cannot reproduce later.

By the end, following along gives you an always-on Pi that serves spoken audio, blocks ads for the house, provides one dashboard URL, works from mobile devices, and can route traffic through a VPN when that makes sense. The chapters are short on purpose: do one job, verify it, move on.

Volume I: Building the homelab

1. **[Foundation: your Pi and your Tailscale network](#)**. Flash a headless 64-bit Linux with key-based SSH, install Docker, and put the Pi, a laptop/desktop, and a phone on one private Tailscale network with MagicDNS, reachable by name from anywhere.
2. **[Streaming audiobooks with Audiobookshelf](#)**. Serve spoken audio from the Pi, streaming to a phone with proper resume tracking. The running example is my own: a language course I ripped from CD so I can study from anywhere without carrying the discs.
3. **[Network-wide ad blocking with Pi-hole](#)**. Run Pi-hole as a DNS server and point your router at it, so every device on your home network is ad-blocked with zero per-device setup.
4. **[Pretty URLs: a reverse proxy and local DNS](#)**. Add Caddy and local DNS so you can type `https://pihole.home` and `https://abs.home` on your tailnet, no ports or IPs, with real HTTPS.
5. **[A one-URL dashboard](#)**. Add Homepage and Portainer, slotted into Caddy at `https://home.home` as a single landing page.

Volume II: VPNs and remote access

6. **[Why a VPN, and Mullvad vs. a standard app like Aura](#)**. What a VPN does, how a mainstream app (I started on Aura) differs from Mullvad, and why Mullvad's portable WireGuard configs are what let a VPN mesh with a homelab, so you can pick the right kind.
7. **[Exit nodes, the Mullvad add-on, and what I run](#)**. The one rule for why homelab access, Pi-hole, and a privacy VPN can't all run at once with two apps, the switch-per-activity setup I live with, what a Tailscale exit node is, and the paid add-on that buys you both at once.

Volume III: Extras

8. **Remoting into machines from a phone.** Reach your Pi’s terminal (Termius) and a full desktop (NoMachine) from your phone or tablet, by Tailscale machine name.
9. **Wiring a phone into Linux machines.** LocalSend and KDE Connect for file transfer and clipboard sharing over Tailscale, plus the USB fallback for bulk photo offload.

Two **appendices** back the parts up: [Appendix A](#) explains every Docker Compose file line by line (kept out of the main chapters so they stay readable), and [Appendix B](#) is a “what should just work” reference, every URL and the one-line check that proves it’s healthy, ideal as a post-reboot smoke test.

The hardware and accounts

- **The always-on server.** I’m running this on a **Raspberry Pi 4 Model B (8 GB RAM)** with a **32 GB microSD card**, in an **Argon ONE M.2 Aluminum case** (passive-plus-fan cooling, with an M.2 slot for a future SSD), on **Ubuntu Server 26.04 LTS (64-bit)**, **headless** (no desktop, just SSH). That’s what I happen to use, and what these steps are written and tested on. You don’t need the exact same kit; anything similar (a Pi 4 or 5 with a few GB of RAM and a 16 GB+ card, on a Debian/Ubuntu base) will follow along. There’s plenty of headroom here, since all the containers, configs, and a ripped audio course together use only a few GB.
- **A laptop/desktop** (Linux assumed) to flash the card and use as a client.
- **A free Tailscale account.** One account, signed into on every device, is the connective tissue for the entire series.
- **A phone.** My examples use an **iPhone**, because that is the phone I have for testing. The same apps have Android builds unless a chapter calls out an iOS-specific limitation, so use the equivalent app for the reader’s platform.
- **Docker** on the Pi, installed once in Chapter 1, reused by every later chapter.
- A Mullvad subscription *or* the Tailscale Mullvad add-on for Chapters 6 and 7 (both optional, ~\$5/month each; Chapter 6 picks the kind of VPN and Chapter 7 picks how it meets your homelab).

Conventions used throughout

- I’m running **Ubuntu Server 26.04 LTS** on the Pi, which is a **Debian/Ubuntu-flavored** Linux, so commands use **apt**, **systemd**, and **docker**.
- Commands prefixed with **sudo** need administrator rights; the rest run as your normal user.
- **The example phone is an iPhone.** I write from the device I can test: Audiobookshelf, Tailscale, Termius, NoMachine, and LocalSend on iOS. That is not an assumption about the reader’s phone. Where the Android path differs, I call it out; otherwise use the same app on the platform at hand.
- **Placeholders, substitute your own.** Wherever you see a Tailscale address like `100.x.y.z`, run `tailscale ip -4` on that machine to get the real one. This guide names the Raspberry Pi homelab on the tailnet; its full DNS name is therefore `homelab.<your-tailnet>.ts.net`,

where `<your-tailnet>` is the suffix shown on the **DNS** page of the Tailscale admin console (something like `tail1a2b3.ts.net`). Substitute it wherever you see `your-tailnet.ts.net`. Likewise, `/home/you/` stands in for your own home directory, and `you@homelab` for your own username on the Pi.

- **Toy example values, swap in yours.** Some things need a concrete number to show, so the guide picks a running example and reuses it. The big ones are the Pi's home (LAN) IP, written as `192.168.1.50`, and the router, written as `192.168.1.1`. *Yours will almost certainly differ* (find the Pi's with `hostname -I`, the router's on the back of the router or in its app). Whenever you see one of these, read it as "e.g." and use your own. The same goes for any bare LAN IP in a command or screenshot.
- **Manual commands first, scripts beside them.** The guide shows the actual files and commands so you can see what is happening. The chapter folders also include numbered helper scripts (`01-...sh`, `02-...sh`, etc.) for the same file-creation and startup steps if you would rather automate the copy/paste. Those scripts are standalone: they write their files inline and do not depend on the `compose/` reference files. Read a script before running it; it is intentionally small enough to audit.

i Why Tailscale is the spine of this series

Every part avoids port forwarding, dynamic DNS, and exposing services to the public internet. Instead, each service binds to the Pi and is reached over an authenticated WireGuard mesh. The result: your homelab is private by default, works on locked-down corporate Wi-Fi, and needs no router configuration. You sign in once per device and the names resolve over the tailnet.

! Security & secrets: read before you publish or push anything

This guide is written to be **safe to commit to a public Git repo** and safe to follow. Two habits make that true, and they're applied throughout:

- **Never put a real password, API token, or key in a file you commit.** Every config below reads its secrets from a local `.env` file that is listed in `.gitignore` and never leaves your machine.
- **Nothing is exposed to the public internet.** Services are reached only over your private tailnet or your home LAN, never by port-forwarding. Remote access to the Pi uses SSH over Tailscale with **key-based authentication** (set up in Chapter 1). Don't open router ports to any of this.

If you fork this guide for your own site, scrub any real tailnet names, IPs, and usernames first, they're mildly identifying. This guide uses placeholders for exactly that reason.

Part I

Volume I: Building the homelab

1 Chapter 1. Foundation: your Raspberry Pi and your Tailscale network

The payoff of this chapter: an always-on Raspberry Pi that’s ready to host everything in this series, running Docker, hardened SSH, and joined to a private Tailscale network, plus a laptop/desktop and phone on that same network, so every machine can reach every other by name, from anywhere, with nothing exposed to the public internet.

This is the groundwork the rest of the book stands on. We don’t install any “app” here, instead we build the platform: a Pi that’s online, secured, and addressable, and a **tailnet** (your private Tailscale network) that stitches your devices together. Every later chapter (Audiobookshelf, Pi-hole, the dashboard, the VPN) plugs into this foundation.

1.1 Why Tailscale is the spine of everything

[Tailscale](#) is a zero-config mesh VPN built on WireGuard. Once installed on your devices and signed into one account, each device gets a stable private address (a `100.x.y.z`) and can reach the others directly, on any network, with no port forwarding and nothing open to the internet.

Why this instead of port forwarding:

- Port forwarding needs a public IP, router configuration, dynamic DNS if your ISP rotates addresses, and it exposes your services to the entire internet.
- A relay like Cloudflare Tunnel avoids port forwarding but proxies all your traffic through a third party.
- Tailscale punches out from both ends and builds a direct, encrypted WireGuard tunnel. Your Pi never accepts an inbound connection from the public internet, only devices on *your* tailnet can talk to it.

That last point is why this guide never tells you to open a router port: the homelab is private by default and works even on locked-down corporate Wi-Fi.

1.2 Prerequisites

- **The always-on server.** I’m using a **Raspberry Pi 4 Model B (8 GB)** with a **32 GB microSD card**, in an **Argon ONE M.2 Aluminum case**, on **Ubuntu Server 26.04 LTS (64-bit)**, **headless**, plus a power supply. That’s what these steps are written and tested on. You can follow along on anything similar (a Pi 4 or 5 with a few GB of RAM and a 16 GB+ card).

- A **laptop/desktop** (Linux assumed; the commands use `apt/systemd`) to flash the card and to use as a client.
- A **phone**. The examples show an iPhone because that is what I have for testing; Android equivalents exist for the apps used here unless a chapter says otherwise.
- A **free Tailscale account**, sign up with Google/GitHub/Microsoft or email at <https://tailscale.com>. One account, signed into on every device, is the only “server” you need.

1.3 Part A: Flash and first boot

1.3.1 Step 1: Flash a headless 64-bit Linux with key-based SSH

Install **Raspberry Pi Imager** on your laptop. I’m running **Ubuntu Server 26.04 LTS (64-bit)**, so in Imager I choose **Other general-purpose OS → Ubuntu → Ubuntu Server 26.04 LTS (64-bit)**. It’s headless, just SSH and a shell, which is all a server needs.

Before writing the card, open Imager’s settings (the gear / **Edit settings**):

- Set a **hostname** (e.g. `homelab`) and your **username**.
- Under **Services**, enable **SSH → “Allow public-key authentication only”**, and paste your laptop’s public key. If you don’t have one yet, generate it first:

```
ssh-keygen -t ed25519          # then paste the contents of
~/.ssh/id_ed25519.pub
```

This bakes key-based login in from first boot and disables password SSH, the single most important hardening step for a box that’s always on.

! Remoting safety for an always-on server

You’ll SSH into this Pi for the rest of the series, so set it up safely once:

- **Key-based auth only, no password login.** The Imager option above handles it; on an already-running Pi you’d `ssh-copy-id you@homelab`, then set `PasswordAuthentication no` in `/etc/ssh/sshd_config` and `sudo systemctl restart ssh`.
- **Reach it over Tailscale, never the public internet.** Once the Pi is on your tailnet (Step 3) you SSH to it as `you@homelab` from anywhere, there is never a reason to port-forward SSH (or anything) on your router. An open port 22 draws constant brute-force traffic; Tailscale sidesteps it entirely.
- **Keep it patched:** `sudo apt update && sudo apt upgrade -y` periodically, or enable `unattended-upgrades`.

Write the card, put it in the Pi, and power on.

1.3.2 Step 2: First boot, then install Docker

Find the Pi on your LAN and SSH in (`ssh you@<pi-lan-ip>`, or `ssh you@homelab.local` if your network supports mDNS). Update it and install Docker, which every later chapter uses to run its services:

```
# Update the base image
sudo apt update && sudo apt upgrade -y

# Docker: the official one-line installer
curl -fsSL https://get.docker.com | sh
sudo usermod -aG docker $USER
newgrp docker          # apply the group change in this shell
```

Why Docker

Every service in this series (Audiobookshelf, Pi-hole, the dashboard, the reverse proxy) ships as a Docker container. Running them in containers keeps their dependencies off your host, makes upgrades a single `docker pull`, and lets each one live in its own tidy folder with a `compose.yaml`. Installing it once here means every later chapter just runs `docker compose up`.

1.4 Part B: Join the tailnet

1.4.1 Step 3: Put the Pi on your tailnet

```
curl -fsSL https://tailscale.com/install.sh | sh
sudo tailscale up
```

`tailscale up` prints a URL, open it, sign in, and the Pi joins your tailnet. Check its address:

```
tailscale ip -4          # something like 100.x.y.z
```

That `100.x.y.z` is how every other device will reach the Pi. You won't have to memorize it, though, MagicDNS (next step) gives it a name.

1.4.2 Step 4: Name the Pi `homelab` and turn on MagicDNS

MagicDNS gives every device on your tailnet a stable name, so you can type `homelab` instead of an IP. In the Tailscale admin console (<https://login.tailscale.com>):

1. **DNS** page → confirm **MagicDNS** is enabled (it's on by default for tailnets created after late 2022; the button reads “Disable MagicDNS” when on). Note your **tailnet name** here, something like `tail1a2b3.ts.net`; wherever this guide writes `your-tailnet.ts.net`, substitute it.
2. **Machines** page → the `⋮` menu on the Pi's row → **Edit machine name** → set it to `homelab`.

Now the Pi answers to `homelab` (and its full name `homelab.your-tailnet.ts.net`) from any device on the tailnet. Because MagicDNS installs your tailnet name as a **search domain**, the short name `homelab` usually works on its own; the full `homelab.your-tailnet.ts.net` is the always-reliable fallback.

1.5 Part C: Add your devices and verify

1.5.1 Step 5: Install Tailscale on your laptop

```
curl -fsSL https://tailscale.com/install.sh | sh
sudo tailscale up          # sign in with the SAME account
```

The tailnet is one shared namespace; signing in with the same account adds the laptop as another peer. From now on you can `ssh you@homelab` from the laptop over the tailnet, anywhere.

1.5.2 Step 6: Install Tailscale on the example phone

Install **Tailscale** from the phone's app store and sign in with the same account. The phone should show `homelab` (and the laptop/desktop) as peers. That's it, the phone is now on the mesh and can reach the Pi by name, on Wi-Fi or cellular.

1.5.3 Step 7: Verify the mesh

From the Pi (or your laptop):

```
tailscale status          # lists every device with a green marker when online
ping homelab              # from another tailnet device, resolves via MagicDNS
ssh homelab               # from your laptop, log straight in by name
```

You should now be able to `ssh homelab` **directly**, no IP and no LAN address, from any device on your tailnet, anywhere. (That works as long as your laptop username matches the one on the Pi; if it differs, use `ssh you@homelab`. The first connection may ask you to confirm the host fingerprint, which is normal.)

If `homelab` resolves, `tailscale status` shows the laptop/desktop and phone online, and `ssh homelab` logs you in, the foundation is done.

💡 A naming note you'll be glad to know later

`nslookup homelab` and `host homelab` can *fail* on macOS even when everything works, those tools bypass the resolver MagicDNS hooks into. Test with `ping homelab` or a browser instead, and keep `homelab.your-tailnet.ts.net` as the guaranteed-resolvable fallback. This trips people up; it's not a broken setup.

1.6 Recap

- **Flashed** a headless 64-bit Linux (Ubuntu Server 26.04 LTS) with key-only SSH (no passwords, no open ports).
- **Installed Docker** on the Pi, the runtime for every later service.
- **Built the tailnet:** the Pi, a laptop/desktop, and a phone all on one private WireGuard mesh, signed into one account.
- **Named the Pi homelab** via MagicDNS, so everything is reachable by name.

The platform is ready. In [Chapter 2](#) we put it to work: ripping my Assimil discs as the working example, then serving audio from the Pi to a phone over the tailnet.

2 Chapter 2. Streaming audiobooks to a phone with Audiobookshelf

The payoff of this chapter: a collection of spoken audio (audiobooks, podcasts, language courses, anything) copied once and served from the Raspberry Pi from Chapter 1, streaming to a phone from anywhere over the tailnet, with proper audiobook-style resume tracking (pause mid-sentence today, resume at the same second tomorrow).

i The concept is general; the example is mine

This part shows how to serve **any** spoken audio from your Pi. The running example I use throughout is my own use case: a language-learning audio course I ripped from CD, because I wanted to study on my lunch break with just the book and none of the discs. Wherever you see the course, picture your own audiobooks; the steps are the same.

In [Chapter 1](#) you built the platform: a Pi named `homelab`, Docker installed, tailnet reachable from a laptop/desktop and phone. This is where it starts doing something useful. **Audiobookshelf** is a self-hosted server for spoken audio. It runs in Docker, scans a folder of MP3s, and gives you a web UI plus iOS and Android apps.

Why an audiobook server and not a music server: music servers (Navidrome, Plex, Jellyfin) track “what song you played last.” Audiobook servers track “second 247 of track 14 of book X” and sync that position across every device. For anything where you pause mid-sentence and resume tomorrow, like a language course or a long audiobook, resume tracking is the feature. Audiobookshelf here is strictly for **spoken** audio: audiobooks, courses, and (later) podcasts. If I add music down the line it gets its own dedicated server, like **Navidrome**, in its own folder, not this one.

In my example the audio starts on CDs, so I need a CD/DVD drive somewhere. Use whatever you’ve got: your laptop or desktop’s built-in drive, or a USB external plugged into the laptop or the Pi. It doesn’t matter which, because the goal is the same either way, the audio ending up in `~/audiobookshelf/media/Audiobooks` **on the Pi**. If your audio is already files on disk, skip the ripping and just copy it into `~/audiobookshelf/media/Audiobooks` on the Pi (over the tailnet with `scp` or `rsync`, or off a USB stick).

2.1 Prerequisites

- You finished [Chapter 1](#): the Pi (`homelab`) is on the tailnet with Docker installed, and at least one client device is on the tailnet too.

- A **CD/DVD drive with the disc inserted**, on whichever machine is handy: a built-in drive, or a USB external plugged into the laptop or the Pi. You'll do the ripping on that machine, then make sure the files land on the Pi.
- The Pi (homelab) reachable over SSH for the server steps.

On the machine that has the drive, check it's seen:

```
lsblk -o NAME,LABEL,FSTYPE,MOUNTPOINT | grep -i -E "rom|cdrom|sr"
```

You should see **sr0**. If it isn't mounted (a headless Pi won't auto-mount it), mount it yourself; this prints the mount point (something like `/media/you/<LABEL>`):

```
udisksctl mount -b /dev/sr0
```

2.2 Part A: Rip and copy the audio

2.2.1 Step 1: Inspect the disc layout

Don't assume the layout. List the mount point (substitute the real label that `udisksctl` just printed):

```
ls /media/$USER/<disc-label>/
```

A language-course data CD like mine typically has three useful things:

1. A **folder of full lessons** with 100 files named `L001-LESSON.mp3` ... `L100-LESSON.mp3`: one complete lesson per file, pre-assembled in playback order.
2. **Per-lesson folders** (`L001-...` ... `L100-...`) that split each lesson into tiny per-sentence files. You can ignore these; the full lessons above are all you need.
3. A **User's Manual** in HTML and PDF.

Plus an `autorun.inf` (a Windows launcher you can ignore on Linux). It's a **data CD**, so this is the fast path: a plain file copy, no audio-CD ripping or re-encoding.

2.2.2 Step 2: Use the one-file-per-lesson folder

Copy the folder with **one complete lesson per file**. That way Audiobookshelf treats the course as a single audiobook with 100 sequential chapters: tap "next chapter" to advance one lesson, and your resume position lives at the lesson level, which is the right unit of progress for a language course.

2.2.3 Step 3: Get the audio onto the Pi

Because it's a data CD, "ripping" is just copying files off the mounted disc. The audiobook library lives at `~/audiobookshelf/media/Audiobooks` on the Pi (we keep the content tidy inside the Audiobookshelf folder, set up in Step 4). Do the copy on whichever machine has the drive.

2.2.3.1 3a. Copy the lessons

If the drive is on the Pi, copy straight into the library:

```
mkdir -p ~/audiobookshelf/media/Audiobooks/"My Course"
rsync -av --no-perms --progress \
    "/media/$USER/<disc-label>/<lessons-folder>/" \
    ~/audiobookshelf/media/Audiobooks/"My Course"/
```

`rsync` over `cp` because if the disc has a read error or you Ctrl-C halfway, re-running resumes where it stopped. `--no-perms` matters: a CD is mounted read-only (files at `r-----`), and the `-a` archive flag would copy those modes to the destination, leaving files you can barely read and directories you can't write into. `--no-perms` uses your `umask` defaults instead, so the result behaves like normal files you created. The L001 ... L100 names are already zero-padded, so Audiobookshelf orders them correctly.

2.2.3.2 3b. If the drive is on your laptop

Rip into any folder on the laptop, then push it to the library path on the Pi over the tailnet (works from anywhere, fastest on the same LAN):

```
rsync -av --no-perms --progress \
    "/media/$USER/<disc-label>/<lessons-folder>/" ~/ "My Course"/
rsync -avz --partial --no-perms --progress \
    ~/ "My Course"/ you@homelab:~/audiobookshelf/media/Audiobooks/"My Course"/
```

`-z` compresses over the wire and `--partial` keeps partial files so an interrupted transfer resumes quickly, useful for a GB-scale copy.

2.2.3.3 3c. Optional: save the manual

To keep the disc's User's Manual alongside the lessons, copy its folder into `~/audiobookshelf/media/Audiobooks/My Course/_manual/` the same way. The leading underscore sorts it to the bottom so Audiobookshelf scans the lessons first.

2.2.3.4 3d. Unmount and eject

When the copy is done, unmount and eject on the machine that had the drive:


```
udisksctl unmount -b /dev/sr0
eject /dev/sr0
```

You only do this once; from now on the files live on the Pi's disk.

2.3 Part B: Run the server

2.3.1 Step 4: Run Audiobookshelf on the Pi

SSH into the Pi if you aren't already there (`ssh you@homelab`); Docker is installed from Chapter 1. We'll run Audiobookshelf as a small **Docker Compose** project, one folder with one `compose.yaml`, the same self-contained pattern every other service in this series uses (so it's one lifecycle to learn and a single file to back up).

```
mkdir -p ~/audiobookshelf/config ~/audiobookshelf/metadata #
media/Audiobooks already exists from Step 3
cd ~/audiobookshelf

cat > compose.yaml <<'EOF'
services:
  audiobookshelf:
    container_name: audiobookshelf
    image: ghcr.io/advplyr/audiobookshelf:latest

    ports:
      - "13378:80"          # host 13378 -> container 80 (host 80 stays free
for Caddy)

    volumes:
      - ./media/Audiobooks:/audiobooks    # add ./media/Podcasts:/podcasts
here later
      - ./config:/config
      - ./metadata:/metadata

    restart: unless-stopped
EOF
```

Then keep the runtime data out of version control and launch:

```
printf 'config/\nmetadata/\nmedia/\n' > .gitignore
docker compose up -d
```

If you are using the chapter scripts instead of typing the file by hand, run these in order from the repo checkout:

```
02-audiobookshelf/scripts/01-make-dirs.sh
02-audiobookshelf/scripts/02-create-compose.sh
02-audiobookshelf/scripts/03-start-audiobookshelf.sh
```

i Why everything lives under ~/audiobookshelf/

The whole stack is **self-contained** in one folder: the `compose.yaml`, the app's `config/` and `metadata/`, and the content itself under `media/`. Because all the volume paths are relative (`./media/Audiobooks`, `./config`), you can move or back up the entire homelab service by copying that one directory. When you add a **Podcasts** library later, it's just another folder, `~/audiobookshelf/media/Podcasts`, mounted at `/podcasts`. (Audiobookshelf is for *spoken* audio; music gets its own server, see the note below.)

What the key settings do:

- **restart:** `unless-stopped` auto-starts the container on boot and after crashes, exactly what you want on an always-on Pi.
- **ports:** `13378:80` maps host port 13378 to the container's port 80 (we avoid host port 80 to leave it free for the reverse proxy in Chapter 4).
- The **volumes** expose your audio under `media/` to the container and persist the database (listening positions, users, settings) next to the compose file, so recreating the container loses nothing. The paths are relative to `~/audiobookshelf/`.

Verify it's running:

```
docker compose ps
curl -I http://localhost:13378      # expect HTTP/1.1 200 OK
```

2.4 Part C: Connect and use it

2.4.1 Step 5: Initial setup in the browser

From any browser on your tailnet, open `http://homelab:13378`. The first-run wizard creates the root user, then drops you at an empty dashboard.

2.4.1.1 5a. Create the root user

This is your admin account. Pick a strong password, Audiobookshelf has no password-reset flow; if you forget it, recovery means editing the SQLite database in `/config` by hand.

2.4.1.2 5b. Add the primary library

Settings (gear) → Libraries → Add Library:

- **Media type: Books.** (This changes the *data model*: Books track one resume position per book in seconds; Podcasts track per-episode played booleans. You want the first for language lessons.)
- **Library name:** *My Course* (specific) or *Audiobooks* (ages better if you'll add more courses). Pick one style and stick to it.
- **Metadata providers:** leave **Open Library** enabled. Even for a homemade rip of a published course it found a clean match for me (cover art and title), where the audiobook-store providers mostly draw blanks on non-commercial audio. You can uncheck the rest to keep scans quick. (For real published audiobooks later, Audible and Google Books are good additions.)
- **Folder:** browse to `/audiobooks`, the *container* path, not the host path. Inside Docker, `~/audiobookshelf/media/Audiobooks` is mounted at `/audiobooks` (from the `volumes:` mapping in the `compose.yaml` in Step 4). If you see `/home/you/Audiobooks` in the picker instead, something's wrong with the volume mount.

Save. Audiobookshelf scans and creates one audiobook with 100 chapters (L001-LESSON.mp3 ... L100-LESSON.mp3) in about 30 seconds.

2.4.1.3 5c. Verify the scan

Open the library, open *My Course*: you should see 100 numbered chapters. Click chapter 1. It should play in the browser. If it does, the server is fully functional.

2.4.2 Step 6: Connect the phone app

1. Install **Audiobookshelf** on the phone. The example here uses the iOS app from the App Store (by advplyr, same as the server); Android has its own build. Tailscale is already installed and signed in from Chapter 1.
2. Open it, tap **Add Server**:
 - **Server address:** `http://homelab:13378` (or `http://100.x.y.z:13378` using the Pi's Tailscale IP).
 - **Username / password:** the root user from Step 5a.
3. Tap your library → your audiobook → the first track. Audio should stream within a second or two, on Wi-Fi or cellular, anywhere.

2.4.3 Step 7: Daily use

- **Playback speed:** 0.5×–3.0×. 0.85× is great for shadowing dialogue you don't fully understand; 1.25× for review passes. (Enable "Remember playback speed per book" in iOS settings, or it resets to 1.0× each track.)
- **Skip & scrub:** ±10s / ±30s buttons; tap-and-hold to scrub.
- **Bookmarks:** drop a named marker at the current position for phrases to revisit.
- **Cross-device progress:** pause in the phone app, open the web UI at home, resume at the same second.

- **Offline downloads:** tap the download icon to cache a book locally, your saving grace if a network is hostile to VPNs. Downloaded audio plays with no network at all.

2.5 Maintenance

Update Audiobookshelf (on the Pi):

```
cd ~/audiobookshelf
docker compose pull && docker compose up -d
```

Because config and metadata live in `~/audiobookshelf/`, recreating the container preserves everything.

Back up listening progress with a cron job on the Pi:

```
0 3 * * * tar czf ~/backups/abs-$(date +%F).tar.gz ~/audiobookshelf/config
~/audiobookshelf/metadata
```

Add more material: drop another folder under `~/audiobookshelf/media/Audiobooks/` (e.g. a Pimsleur course) and click “Scan Library.” Each top-level folder becomes its own audiobook.

2.6 Troubleshooting

The phone app says “Cannot connect.” From the phone’s browser, visit `http://homelab:13378`. If the browser also fails, Tailscale isn’t up, open the Tailscale app and confirm both devices show green dots. If Safari works but the app fails, you typed the wrong port (13378).

Tracks play in the wrong order. Audiobookshelf sorts by filename. Rename so they sort correctly (Lesson 01.mp3, not Lesson 1.mp3), then “Re-scan” in the web UI.

Container won’t restart after a reboot. `cd ~/audiobookshelf && docker compose ps`; if it’s down, `docker compose up -d`. If `restart: unless-stopped` isn’t taking effect, check `systemctl status docker`, Docker itself may not be enabled at boot.

rsync reports “Permission denied” on every mkdir. Check the destination mode: `ls -ld /path`. A `dr-x-----` directory is missing its write bit; fix with `chmod u+w /path` (or `chmod 755`) and re-run.

Work Wi-Fi blocks Tailscale. Tailscale falls back to DERP relays automatically (slower but works). Last resort: download lessons in the phone app before leaving home, cached audio plays fully offline.

2.7 Recap

- **Copied** the audio off the disc straight onto the Pi with `rsync --no-perms`.
- **Ran** Audiobookshelf on the Pi as a small `docker compose` project.
- **Connected** the phone app to `http://homelab:13378`.

Your course is now permanently available from any device on your tailnet, with proper progress tracking and no subscription. Next, in [Chapter 3](#), we give the Pi a second job: blocking ads for every device on your home network with Pi-hole.

3 Chapter 3. Network-wide ad blocking at home with Pi-hole

The payoff of this chapter: every device on your home network, laptops, phones, the smart TV, game consoles, IoT gadgets, has ads and trackers stripped out *before they ever load*, with zero per-device configuration.

Across Chapters 1–2 you put an always-on Raspberry Pi (`homelab`) on your tailnet and gave it its first job, Audiobookshelf. Pi-hole is the natural second tenant: it’s a **DNS sinkhole**, a DNS server that answers “no such host” for domains known to serve ads, trackers, and malware, and forwards everything else to a real upstream resolver. Because it works at the DNS layer, it blocks ads in *every* app and browser at once, including places a browser extension can’t reach (smart TVs, mobile apps, in-game ads).

This chapter is about the home network, not remote access yet. Pi-hole becomes the DNS server for the house. Set it once in the router and every device on your Wi-Fi inherits it automatically: no per-device tweaking, no VPN decisions, just clean DNS under your roof.

3.1 What you are building (and why this design)

The plan has two halves:

1. **Run Pi-hole on the Pi** (in Docker, alongside Audiobookshelf).
2. **Point your home network at it** by setting Pi-hole as the DNS server in your router, so every device that joins your Wi-Fi automatically uses it.

That router-level step is the payoff: configure one setting in one place, and every current and future device on your home network is covered without touching any of them individually.

3.2 Part A: Run Pi-hole on the Pi

3.2.1 Step 1: Free up port 53 on the Pi

This is the one host-level snag, and it bites almost everyone, so we do it first. Pi-hole needs to listen on **port 53** (the DNS port). But Raspberry Pi OS (like Debian/Ubuntu) ships `systemd-resolved`, which runs a small “stub” DNS listener bound to `127.0.0.53:53`. Two programs can’t own the same port, so Pi-hole will fail to start until you evict the stub.

The surgical fix is to disable *only* the stub listener while leaving `systemd-resolved` running to handle the Pi’s own name resolution:

```
# 1. Disable the stub listener via a drop-in (survives package upgrades)
sudo mkdir -p /etc/systemd/resolved.conf.d
printf '[Resolve]\nDNSStubListener=no\n' \
| sudo tee /etc/systemd/resolved.conf.d/no-stub.conf

# 2. Point the Pi's own resolver at resolved's real file, not the dead stub
sudo ln -sf /run/systemd/resolve/resolv.conf /etc/resolv.conf

# 3. Apply the change
sudo systemctl restart systemd-resolved

# 4. Confirm port 53 is now free (this should print nothing)
sudo ss -tulpn | grep ':53'
```

💡 Why a drop-in file instead of editing the main config

Dropping a file into `/etc/systemd/resolved.conf.d/` rather than editing the main `resolved.conf` keeps your change isolated and reversible: undoing it is `sudo rm /etc/systemd/resolved.conf.d/no-stub.conf` followed by a restart, with no risk of clobbering a setting a future package update wants to change in the main `resolved.conf`.

If step 4 prints a line mentioning `systemd-resolve`, the stub is still bound, re-check that the drop-in saved correctly and that you restarted the service.

3.2.2 Step 2: Give the Pi a fixed LAN IP

Pi-hole's address must never change, because in a moment your whole network will point its DNS at it, and a Pi-hole whose IP drifts is a house-wide DNS outage. So before configuring anything, pin the Pi to one LAN address.

In your **router's admin page**, find the DHCP settings and add a **reservation** (sometimes called a *static lease*): tell the router to always hand the Pi the same IP. Match it by the Pi's MAC address, or pick the Pi out of the list of connected devices. Note the address it gets (for example 192.168.1.50); if you need the Pi's current IP, run `hostname -I` on the Pi. You'll use this address as `PIHOLE_IP` in the next step.

No reservation option on your router? You can set a static IP on the Pi itself instead, but a router reservation is simpler and the usual way.

3.2.3 Step 3: Run Pi-hole in Docker

You already have a Docker workflow from Chapter 1, so we'll keep Pi-hole in the same world rather than installing it natively. That means one lifecycle to learn (`docker ...`), declarative config you can back up as a file, and no new system packages on the host.

We'll keep this to one clean, self-contained Compose project: a single folder holding one `compose.yaml`, one `.env` for your settings and secrets, and one `.gitignore`. You can read it, back it up, and version-control it without ever exposing a password.

1. Make the folder.

```
mkdir -p ~/pihole
cd ~/pihole
```

2. Create the `.env`. This is the only place your timezone, Pi IP, and admin password live:

```
cat > .env <<'EOF'
TZ=America/Denver
PIHOLE_PASSWORD=replace-this-with-a-strong-password
PIHOLE_IP=192.168.1.50
EOF
chmod 600 .env      # readable only by you
```

Set `TZ` to your timezone and `PIHOLE_IP` to the address you reserved in Step 2. For the password, type a strong one, or generate one and drop it straight in:

```
# optional: replace the placeholder with a random password
sed -i "s|^PIHOLE_PASSWORD=.*|PIHOLE_PASSWORD=$(openssl rand -base64 18)|"
.env
```

💡 Why `PIHOLE_IP`, and why bind the ports to it

The compose file below publishes Pi-hole's ports on `${PIHOLE_IP}` specifically (e.g. `192.168.1.50:53`) rather than on every interface. Serving on the one LAN address in use keeps Pi-hole listening where it needs to and nowhere else. Make `PIHOLE_IP` the **static** or **DHCP-reserved** address you reserved in Step 2, so it never changes underneath you.

3. Create `compose.yaml`.

```
cat > compose.yaml <<'EOF'
services:
  pihole:
    container_name: pihole
    image: pihole/pihole:latest

    ports:
      - "${PIHOLE_IP}:53:53/tcp"
      - "${PIHOLE_IP}:53:53/udp"
      - "${PIHOLE_IP}:80:80/tcp"

    environment:
```



```

    TZ: "${TZ}"
    FTLCNF_webserver_api_password: "${PIHOLE_PASSWORD:?set
PIHOLE_PASSWORD in .env}"
    FTLCNF_dns_upstreams: "1.1.1.1;1.0.0.1"
    FTLCNF_dns_listeningMode: "ALL"

    volumes:
      - ./etc-pihole:/etc/pihole

    restart: unless-stopped
EOF

```

There is **no password in this file**, only a `${PIHOLE_PASSWORD}` reference that Compose fills in from `.env` at startup. The `${PIHOLE_PASSWORD:?...}` form makes Compose stop with a clear error if that variable is missing, so you can't accidentally launch with a blank password.

4. Create the `.gitignore` so the secret and the runtime data can never be committed:

```

cat > .gitignore <<'EOF'
.env
etc-pihole/
EOF

```

! The pattern, in one line

If it's a password, token, or key, it goes in `.env`, and `.env` goes in `.gitignore`. The `compose.yaml` holds only `${VARIABLE}` references; the real values live in `.env`, which never leaves your Pi. Every later service in this series uses the exact same pattern.

5. Start it.

```

docker compose up -d
docker compose logs --tail 20      # watch it come up

```

Script version, same order as the manual steps:

```

03-pihole/scripts/01-free-port-53.sh
03-pihole/scripts/02-create-env.sh
03-pihole/scripts/03-create-compose.sh
03-pihole/scripts/04-create-gitignore.sh
03-pihole/scripts/05-start-pihole.sh

```

i Why port 80 now, and why it moves to 8081 later

A web UI conventionally lives on **port 80**, the default HTTP port, which is why Pi-hole's admin sits there for now. We're only borrowing it. In [Chapter 4](#) we add a reverse proxy that has to **own** port 80, because that's the port a browser connects to when you type a bare name like `pihole.home` with no `:port`. Only one program can hold a host port, so the front door (Caddy) gets 80 and Pi-hole's web UI moves to **8081** then. The number mostly disappears once the URL is `pihole.home` anyway. Why 8081 and not 8080? Counter-intuitively, **8080 is the most common alternate HTTP port** (dev servers, lots of containers grab it), so it's the one most likely to be taken; 8081 is just one past it. Until then, make sure nothing else holds port 80: Audiobookshelf (Chapter 2) uses 13378 so it's clear, but `sudo ss -tlnp | grep ':80'` confirms it before `docker compose up`.

3.2.4 What the key settings do

- **FTLCONF_* environment variables** are Pi-hole v6's configuration mechanism. Each one maps to a key in Pi-hole's single config file (`/etc/pihole/pihole.toml`): `FTLCONF_<section>_<key>` → `[section] key`. Anything you set this way becomes **read-only in the web UI**, it's re-applied from the environment on every container start, which is exactly what you want for reproducible, file-defined config.
- **FTLCONF_webserver_api_password** is the v6 replacement for the old `WEBPASSWORD` variable. If you omit it, Pi-hole generates a random password and prints it to the container log on first boot.
- **FTLCONF_dns_upstreams** is who Pi-hole asks when a domain *isn't* blocked. Cloudflare (1.1.1.1) is fast; Quad9 (9.9.9.9) adds its own malware filtering. Pick one.
- **FTLCONF_dns_listeningMode: "ALL"** tells Pi-hole to answer queries arriving on any interface. In Docker's bridge network, your LAN clients' queries reach Pi-hole *through* the Docker gateway rather than appearing to come from your local subnet, so the stricter "local only" mode would drop them. **ALL** is the documented working choice for a Dockerized Pi-hole.

i Pi-hole v6 is meaningfully different from the v5 you'll find in old guides

Pi-hole **v6** (early 2025) collapsed several moving parts into one. There's no more `lighttpd` web server and no PHP, the `pihole-FTL` process now serves the DNS resolver, the web dashboard, *and* a REST API itself. Configuration moved from a pile of files (`setupVars.conf`, `pihole-FTL.conf`, `dnsmasq` snippets) to a single `/etc/pihole/pihole.toml`. The commands changed too: `whitelist/blacklist` became `pihole allow/pihole deny`, and the password command is now `pihole setpassword`. If a tutorial tells you to edit `setupVars.conf`, it predates v6, ignore it.

3.2.5 Step 4: Set the admin password and log in

The web UI lives at `http://<pi-ip>/admin`. On your home network use the Pi's LAN IP (e.g. `http://192.168.1.50/admin`); find it with `hostname -I` on the Pi.

Because the password comes from `PIHOLE_PASSWORD` in your `.env`, it's already configured, log in with whatever value you put there. To change it later, edit `.env` and recreate the container:

```
# edit ~/pihole/.env, then:  
cd ~/pihole && docker compose up -d --force-recreate
```

⚠ Don't fight your own `.env`

Because the password is pinned by the `FTLCONF_webserver_api_password` environment variable, running `pihole setpassword` inside the container won't stick, the value from `.env` is re-applied on every restart and silently wins. With this setup, the `.env` file is the single source of truth for the password; change it there, nowhere else.

3.3 Part B: Point your home network at it

3.3.1 Step 5: Point your home network at Pi-hole

This is the centerpiece. You want every device in the house to use Pi-hole as its DNS server without configuring each one. The usual way is to set it **once at your router**, because the router hands out DNS settings to every device via DHCP when they join the Wi-Fi.

1. Use the Pi's reserved LAN IP from Step 2 (e.g. `192.168.1.50`). Every device will be told to use that address for DNS.
2. In the router's settings, find the **DNS server** field, usually under *DHCP*, *LAN*, or *Internet* settings, and set the **primary DNS** to the Pi's IP (e.g. `192.168.1.50`).
3. **Leave the secondary DNS blank if you can.** This is counterintuitive: a secondary DNS (like `8.8.8.8`) feels like a safety net, but devices freely use *either* server, so half your queries would skip Pi-hole and you'd see ads "randomly." A single DNS entry forces everything through Pi-hole.
4. Renew leases so devices pick up the change: reboot the router, or toggle Wi-Fi off/on on each device (or just wait, leases renew on their own).

i What if your ISP router won't let you change DNS?

Some locked-down ISP gateways don't expose the DNS field. Two fallbacks: (a) set DNS **per-device** (each phone/laptop's Wi-Fi settings let you set a manual DNS, more tedious but works), or (b) let Pi-hole run your network's **DHCP** instead of the router (an option in Pi-hole's settings; you'd disable the router's DHCP first). The router approach above is by far the simplest when it's available.

⚠ Don't expose port 53 to the internet

Whatever you do, never port-forward port 53 from your router to the Pi. An open, internet-facing DNS resolver gets abused for DNS-amplification attacks within hours. Pi-hole answering your *LAN* is exactly right; Pi-hole answering the *whole internet* is a problem. Keep it on your

home network.

3.4 Part C: Add blocklists and verify

3.4.1 Step 6: Add good blocklists

A fresh Pi-hole v6 already subscribes to **StevenBlack's unified hosts list**, which is a strong baseline on its own. You can verify it under **Lists** in the UI. To add more coverage:

1. In the web UI, go to **Lists** (sometimes shown as “Adlists” or “Subscribed Lists”).
2. Paste a list URL and click **Add**.
3. Crucially, go to **Tools** → **Update Gravity** (or run `docker exec pihole pihole -g`). Blocklists do *nothing* until “gravity”, Pi-hole's compiled blocklist database, is rebuilt.

Two well-maintained additions that block aggressively without breaking sites. Each URL is in its own block below so it copies cleanly, paste one at a time into **Lists** → **Add**:

OISD Big, a popular all-in-one list that's deliberately conservative about breaking sites:

```
https://big.oisd.nl
```

HaGeZi Multi PRO, actively maintained and tiered (“PRO” is the balanced middle tier). The raw URL is long; copy the whole single line:

```
https://raw.githubusercontent.com/hagezi/dns-blocklists/main/domains/pro.txt
```

Don't stack a dozen lists

More lists is not more better. Two or three quality lists (StevenBlack + OISD *or* HaGeZi) outperform a sprawling pile of overlapping ones and cause far fewer “this website is broken” surprises. Every time you add or remove a list, run **Update Gravity** or nothing changes.

3.4.2 Step 7: Verify the whole thing

Run through these in order; each isolates a different layer:

```
# Query Pi-hole at the Pi's LAN IP (e.g. 192.168.1.50), the address it's
bound to.
# 1. On the Pi itself:
dig +short google.com @192.168.1.50          # returns IPs = resolver works
dig +short doubleclick.net @192.168.1.50     # returns 0.0.0.0 = blocking
works

# 2. From another device on the LAN (run this from your laptop at home):
```

```
dig +short doubleclick.net @192.168.1.50    # 0.0.0.0 = LAN clients can use
it
```

i Why not @127.0.0.1?

You might reach for `dig @127.0.0.1` on the Pi, but here it returns “connection refused”, and that’s correct, not a fault. In this chapter Pi-hole is published on the **LAN IP only** (`{PIHOLE_IP}:53`), not on localhost, and Step 1 disabled the `systemd-resolved` stub that used to answer on `127.0.0.53`. So query the Pi by its LAN address instead. (In [Chapter 4](#) you switch Pi-hole to listen on all interfaces, and `@127.0.0.1` starts working too.)

Then the real-world test: on a device connected to your home Wi-Fi, open a normally ad-heavy site or app, then watch the **Query Log** at your Pi’s admin page (e.g. `http://192.168.1.50/admin`) from your laptop. You should see that device’s LAN IP making queries, with ad domains marked blocked.

If a device’s queries don’t appear at all, it’s still using its old DNS, renew its DHCP lease (toggle Wi-Fi off/on) so it picks up the router’s new DNS setting.

3.5 Troubleshooting

Pi-hole container won’t start / “port 53 already in use.” The `systemd-resolved` stub is still bound. Re-run Step 1 and verify with `sudo ss -tulpn | grep ':53'`. You should see only Docker/Pi-hole, never `systemd-resolve`.

Some devices still show ads after the router DNS change. Either they haven’t renewed their DHCP lease (toggle Wi-Fi off/on), or your router has a *secondary* DNS set that lets them bypass Pi-hole. Clear the secondary entry (Step 5.3).

A device hard-codes its own DNS and ignores the router. Some devices (notably certain Android phones, Chromecasts, and anything using DNS-over-HTTPS to 8.8.8.8) bypass your router’s DNS entirely. You can’t fix those from the router; either change DNS on the device itself or block their hard-coded resolvers, an advanced topic, but worth knowing it’s *them*, not a broken Pi-hole.

A specific site is broken. A blocklist is over-matching. Find the blocked domain in the **Query Log**, click it, and **Allow** it. Then the site works while everything else stays blocked.

You lost the admin password. It’s in your `.env` file, `cat ~/pihole/.env` and read `PIHOLE_PASSWORD`. (This is the upside of keeping it in `.env`: it’s recoverable from a file only you can read, not lost in a container.)

The Pi itself can’t resolve names after Step 1. Check that `/etc/resolv.conf` is the symlink from Step 1.2 and that your upstream (`FTLCONF_dns_upstreams`) is a real public resolver, not the router, if the router points back at Pi-hole (that’s a resolution loop).

3.6 Recap

- **Free port 53** by disabling the `systemd-resolved` stub listener.
- **Run Pi-hole** in Docker Compose next to Audiobookshelf.
- **Point your home network at it** by setting the Pi as the router's DNS server (single entry, no secondary), one setting covers every device.
- **Add one or two quality blocklists** and run **Update Gravity**.

Your homelab now does two jobs on your home network. In [Chapter 4](#) we stop typing IP addresses and ports to reach these services, a small reverse proxy plus local DNS gives Pi-hole and Audiobookshelf clean names like `https://pihole.home` and `https://abs.home` that work from anywhere on your tailnet.

4 Chapter 4. Pretty URLs: a reverse proxy + local DNS

The payoff of this chapter: type a bare `pihole.home` or `abs.home` in any browser on your tailnet (at home or on cellular) and land on the right service over HTTPS, no port numbers, no IP addresses, no scheme to type. And a pattern you'll reuse for every service you add from here on.

Right now the setup works, but the URLs are ugly: `http://192.168.1.50/admin` for Pi-hole, `http://homelab:13378` for Audiobookshelf, an IP here, a port there. This chapter fixes that once and gives every service one clean name. When we add the dashboard in [Chapter 5](#), it drops into this same pattern.

! Why not `pihole.com`?

A `.com` (or any public) name belongs to whoever registered it on the internet, you can't point `pihole.com` at your Pi any more than you can point `google.com` at it. We use a **private** suffix, `.home`, that exists only on *your* network.

4.1 How this works: two jobs, two tools

A working URL needs two separate things:

1. **The name has to resolve to your Pi.** `pihole.home` must turn into the Pi's address, that's **DNS** (we add records to Pi-hole and let Tailscale carry them to every device).
2. **Traffic hitting the Pi has to reach the right service.** `pihole.home` and `abs.home` both arrive on the Pi's web ports (80/443), but Pi-hole and Audiobookshelf listen on *different* ports. Something must read the requested name and forward to the correct backend, a **reverse proxy** (we'll use **Caddy**, the simplest to configure). Caddy also gives each name an HTTPS certificate from its own private CA, so the URLs get a real padlock.

<code>pihole.home</code>	DNS	the Pi (443)	Caddy routes by name	<code>pihole:80</code>
<code>abs.home</code>	DNS	the Pi (443)	Caddy routes by name	<code>audiobookshelf:80</code>

4.2 Part A: Wire the services onto one network

4.2.1 Step 1: Create the shared Docker network and put Audiobookshelf on it

Caddy reaches each backend by **container name** (pihole, audiobookshelf), which only works if they share a Docker network. Create it:

```
docker network create homelab
```

Both Audiobookshelf (Chapter 2) and Pi-hole (Chapter 3) are **Compose projects**, so we attach each to the network the *declarative* way, by adding it to the project's `compose.yaml` (see the callout for why that's the only attachment that sticks). Start with Audiobookshelf: add the **networks** blocks to `~/audiobookshelf/compose.yaml`:

```
cat > ~/audiobookshelf/compose.yaml <<'EOF'
services:
  audiobookshelf:
    container_name: audiobookshelf
    image: ghcr.io/advplyr/audiobookshelf:latest

    ports:
      - "13378:80"

    volumes:
      - ./media/Audiobooks:/audiobooks
      - ./config:/config
      - ./metadata:/metadata

    networks:
      - homelab          # NEW: join the shared homelab network

    restart: unless-stopped

networks:
  homelab:
    external: true      # NEW: the shared network created above
EOF
( cd ~/audiobookshelf && docker compose up -d --force-recreate )
```

Any container on the `homelab` network can now reach the others by name, no IP addresses needed between containers.

 **Attach Compose services in the Compose file, not with `docker network connect`**

It's tempting to run `docker network connect homelab audiobookshelf` instead, but it won't survive. Any time you recreate the container (`docker compose up --force-recreate`, or an update), a **recreate rebuilds it with only the networks declared in its**

`compose.yaml`, a network you attached by hand is silently dropped, and Caddy would then fail to resolve the backend. So for every Compose service we add the `homelab` network *inside* the Compose file, where a `recreate` always re-attaches it. Pi-hole gets the same treatment in Step 2.

4.2.2 Step 2: Free port 80 and extend Pi-hole onto the tailnet

Three changes to Pi-hole's `~/pihole/compose.yaml`, all needed before Caddy works:

- **Free host port 80.** Caddy becomes the single front door on port 80, because that's the port a browser hits for a bare URL like `pihole.home`. The backends are reached *through* Caddy by container name (`pihole:80`) on the internal network, so their **host** port no longer matters, which is why we can move Pi-hole's web UI off 80 to 8081.
- **Listen on the tailnet, not just the LAN.** In Chapter 3 you bound Pi-hole to your LAN IP because it was a home-only service. To make `.home` names (and Pi-hole's own ad-blocking) work *everywhere*, Pi-hole's DNS must also answer on the Pi's Tailscale interface, so drop the `${PIHOLE_IP}:` prefix and let it listen on all interfaces.
- **Join the homelab network declaratively** (per the Step 1 callout) so Caddy can reach it as `pihole:80`, and so the `recreate` below, and every future `docker compose up`, re-attaches it automatically.

Replace `~/pihole/compose.yaml` with this updated version (the `ports:` and new `networks:` blocks are the only changes from Chapter 3):

```
cat > ~/pihole/compose.yaml <<'EOF'
services:
  pihole:
    container_name: pihole
    image: pihole/pihole:latest

    ports:
      - "53:53/tcp"          # all interfaces (LAN + tailnet); was
${PIHOLE_IP}:53:53
      - "53:53/udp"
      - "8081:80/tcp"        # web UI moved off 80, freeing it for Caddy;
was ${PIHOLE_IP}:80:80

    environment:
      TZ: "${TZ}"
      FTLCONF_webserver_api_password: "${PIHOLE_PASSWORD:?set
PIHOLE_PASSWORD in .env}"
      FTLCONF_dns_upstreams: "1.1.1.1;1.0.0.1"
      FTLCONF_dns_listeningMode: "ALL"

    volumes:
      - ./etc-pihole:/etc/pihole
```

```

    networks:
      - homelab          # NEW: declarative attachment, survives every
recreate

    restart: unless-stopped

networks:
  homelab:
    external: true      # NEW: the shared network you created in Step 1
EOF

```

Then recreate the container so the new ports and network take effect:

```
cd ~/pihole && docker compose up -d --force-recreate
```

Script version for Part A:

```

04-pretty-urls/scripts/01-create-homelab-network.sh
04-pretty-urls/scripts/02-create-audiobookshelf-network-compose.sh
04-pretty-urls/scripts/03-create-pihole-network-compose.sh
04-pretty-urls/scripts/04-recreate-networked-services.sh

```

i Is listening on all interfaces safe?

Yes, in this setup. The only networks that can reach the Pi are your home LAN (behind the router) and your tailnet (authenticated WireGuard). The one unbreakable rule, same as Chapter 3: **never port-forward port 53 (or 80) from your router**. As long as you don't, "all interfaces" just means "my LAN and my tailnet," which is exactly who should be able to use it.

Pi-hole's admin UI now lives at <http://192.168.1.50:8081/admin> (your Pi's LAN IP), but in a moment you'll reach it as <https://pihole.home> instead.

4.3 Part B: Deploy Caddy and DNS

4.3.1 Step 3: Deploy Caddy

Caddy is its own small stack:

```
mkdir -p ~/caddy && cd ~/caddy
```

Create `~/caddy/Caddyfile`, the entire routing table:

```

cat > Caddyfile <<'EOF'
# `tls internal` makes Caddy issue each .home name a certificate from its
OWN
# built-in certificate authority (no public CA issues certs for a private
TLD).
# Caddy then serves HTTPS on 443 and auto-redirects http:// -> https://.
Once you
# trust Caddy's local CA on your devices (Step 7), a bare `pihole.home`
typed in
# the browser opens over HTTPS: no scheme, no warning, real padlock.

pihole.home {
    tls internal
    redir / /admin 302          # land on the admin page directly
    reverse_proxy pihole:80
}

abs.home {
    tls internal
    reverse_proxy audiobookshelf:80
}
EOF

```

Create ~/caddy/compose.yaml:

```

cat > compose.yaml <<'EOF'
services:
  caddy:
    container_name: caddy
    image: caddy:latest
    ports:
      - "80:80"                # the single front door (redirects to 443)
      - "443:443"              # HTTPS, on all interfaces (LAN + tailnet)
      - "443:443/udp"          # HTTP/3
    volumes:
      - ./Caddyfile:/etc/caddy/Caddyfile:ro
      - caddy_data:/data
      - caddy_config:/config
    networks:
      - homelab
    restart: unless-stopped

networks:
  homelab:
    external: true            # the shared network from Step 1

volumes:

```

```
caddy_data:
caddy_config:
EOF
```

Start it:

```
docker compose up -d
docker compose logs --tail 20
```

Script version for the Caddy files and startup:

```
04-pretty-urls/scripts/05-create-caddyfile.sh
04-pretty-urls/scripts/06-create-caddy-compose.sh
04-pretty-urls/scripts/07-start-caddy.sh
```

Caddy listens on `0.0.0.0:80` and `0.0.0.0:443`, all interfaces, including the Pi's Tailscale `100.x.y.z`, and reaches each backend by container name over the `homelab` network, so it ignores host port mappings entirely. That's why moving Pi-hole's `host` web port to 8081 doesn't bother Caddy: it talks to `pihole:80` inside the network. On first start, `tls internal` makes Caddy generate a local certificate authority and sign a cert for each `.home` name; you'll trust that CA on your devices in Step 7.

4.3.2 Step 4: Add the `.home` names to Pi-hole's DNS

Make the names resolve to the Pi. In the Pi-hole admin UI (e.g. `http://192.168.1.50:8081/admin` for now), go to **Settings** → **Local DNS Records** and add one **A record** per name, all pointing at the Pi's **Tailscale IP** (`tailscale ip -4` on the Pi):

Domain	IP
<code>pihole.home</code>	<code>100.x.y.z</code>
<code>abs.home</code>	<code>100.x.y.z</code>

Or let the helper add both records:

```
04-pretty-urls/scripts/08-add-local-dns-records.sh
```

💡 Why the *Tailscale* IP, not the LAN IP

A DNS record holds one address. Pointing these at the Pi's `100.x.y.z` tailnet address makes them reachable from *anywhere* a device is on your tailnet, home Wi-Fi or cellular alike. Caddy is listening on that address (Step 3), so the traffic lands. (This assumes Tailscale is running on the device, which is the premise of the whole series.)

4.3.3 Step 5: Make Pi-hole your tailnet's DNS

For *.home to resolve on every device, your devices must ask **Pi-hole** for DNS. In the Tailscale admin console (<https://login.tailscale.com>), **DNS** page:

1. **Add nameserver** → **Custom**, enter the Pi's Tailscale IP (100.x.y.z), save.
2. Turn on **Override local DNS**. Now every tailnet device resolves through Pi-hole, so *.home names work, *and* you get Pi-hole ad-blocking on every device, even on cellular.

i Want only .home through Pi-hole, not all your DNS?

Use **split DNS** instead: in **Add nameserver** → **Custom**, enable **Restrict to domain** and set the domain to **home**. Then only *.home lookups go to Pi-hole and everything else uses each device's normal DNS. You lose tailnet-wide ad-blocking but keep the pretty URLs. ([Chapter 7](#) discusses how this DNS choice interacts with VPNs when you're away from home.)

4.4 Part C: Use it and trust the cert

4.4.1 Step 6: Try it

From your laptop or phone, anywhere on the tailnet:

```
https://pihole.home    # Pi-hole admin (Caddy redirects to /admin)
https://abs.home       # Audiobookshelf
```

No ports, no IPs. To confirm resolution, `ping pihole.home` should answer from the Pi's 100.x.y.z. Until you do Step 7, your browser will warn that the certificate isn't trusted (it's Caddy's private CA), that's expected, and Step 7 removes the warning for good.

4.4.2 Step 7: Trust Caddy's certificate (so bare names *just work*)

Without this step, typing a bare `pihole.home` is frustrating: browsers try **https:// first**, and because they don't yet trust Caddy's private CA, they throw a "not secure" warning (or fall back to treating the name as a web search). The fix is to install **Caddy's root certificate** on each device, once. After that, the browser's automatic HTTPS upgrade lands on a *trusted* cert and bare `home.home` / `pihole.home` / `abs.home` open instantly, padlock and all.

First, export the root certificate from the Caddy container (on the Pi):

```
docker exec caddy cat /data/caddy/pki/authorities/local/root.crt \
> ~/caddy/caddy-root-ca.crt
```

Now get that file onto each device. Since every device is already on your tailnet, use **Taildrop**, Tailscale's built-in file send between your own machines, no AirDrop, no email, no cable:

```
# from the Pi (or any machine that has the file), send to a peer by its
name:
tailscale file cp ~/caddy/caddy-root-ca.crt <laptop-name>:
tailscale file cp ~/caddy/caddy-root-ca.crt <iphone-name>:
```

Find the exact peer names with `tailscale status`. (If the Pi answers `Access denied: file access denied`, run `sudo tailscale set --operator=$USER` once so the Tailscale CLI works without `sudo`, then retry, or just send from a machine where it already works.)

Then install `caddy-root-ca.crt` on each device:

- **Linux laptop (system-wide):** receive the Taildrop file (it lands in your Downloads, or run `tailscale file get .`), then:

```
sudo cp caddy-root-ca.crt
/usr/local/share/ca-certificates/caddy-root-ca.crt
sudo update-ca-certificates
```

Firefox keeps its **own** trust store: *Settings* → *Privacy & Security* → *Certificates* → **View Certificates** → *Authorities* → **Import** the file and tick “Trust this CA to identify websites.”

- **iPhone:** open the **Tailscale** app, tap the received `caddy-root-ca.crt`, and **Save to Files**. Open it from the **Files** app → iOS says “Profile Downloaded” → **Settings** → **General** → **VPN & Device Management** → **Install** the profile. Then (this second step is the one people miss) **Settings** → **General** → **About** → **Certificate Trust Settings** and toggle **on** “Enable Full Trust” for the Caddy Local Authority. (AirDrop or email work too, but Taildrop keeps it on the tailnet.)

Now reload `https://home.home`, the warning is gone, and typing the bare name works.

i Why a private CA instead of a real certificate

A real, publicly-trusted certificate requires a domain under the reader’s control and a public CA (Let’s Encrypt) that can verify it, which can’t be done for a private `.home` name. Caddy’s internal CA sidesteps that entirely: it mints certificates locally and you tell *your* devices to trust *that* CA. The trade-off is the one-time install above. (If you ever buy a real domain, point a subdomain at the Pi and drop `tls internal`, Caddy will fetch a real certificate automatically, and you can remove the root cert from your devices.)

4.5 The pattern you’ll reuse

Every service from here on gets a pretty URL the same way, three small steps:

1. Put its container on the `homelab` network.
2. Add a block to `~/caddy/Caddyfile`:
`name.home { tls internal\n reverse_proxy container:PORT }.`

3. Add a `name.home → 100.x.y.z` record in Pi-hole, then `docker compose restart caddy` (in `~/caddy`).

Because you’ve already trusted Caddy’s CA, the new name gets a trusted cert automatically, no per-service certificate work. [Chapter 5](#) does exactly this for the dashboard, giving it `https://home.home`.

4.6 Caveats

- **One-time CA trust per device.** The padlock only appears on devices where you’ve installed Caddy’s root cert (Step 7). A brand-new device shows the warning until you do. This is the cost of a private TLD; it’s a one-time step.
- **Exit nodes break .home.** A Mullvad exit node routes DNS through Mullvad and bypasses Pi-hole, so `*.home` won’t resolve while one is active. This tension is the subject of [Chapter 7](#); switch the exit node off (or use your home Pi as the exit node) to use the names.

4.7 Recap

- **Shared homelab network** so Caddy can reach services by name.
- **Pi-hole** moved its web UI to 8081 and now listens on all interfaces (LAN + tailnet).
- **Caddy** is the single front door on ports 80/443, serving HTTPS with its own internal CA and routing `pihole.home` and `abs.home` to the right containers.
- **Pi-hole DNS records + the Tailscale Override** make those names resolve on every device, everywhere.
- **Trusting Caddy’s root CA** on each device makes bare names open over HTTPS with no warning and no scheme.

You now have clean, portable URLs, and a repeatable recipe for every service to come. Next, [Chapter 5](#) adds a dashboard that ties them all together behind `https://home.home`.

5 Chapter 5. A one-URL dashboard at `https://home.home`

The payoff of this chapter: open any browser on your tailnet, type `https://home.home`, and land on a clean dashboard that shows your services at a glance and links into a full Docker management console, the single front door that ties everything together.

Chapter 4 gave individual services clean names (`https://pihole.home`, `https://abs.home`). This chapter adds the page I want to open first: one dashboard for status and links, plus a real Docker management GUI when I need to touch containers directly. It reuses the reverse-proxy pattern you just built instead of introducing another way to expose things.

We'll install two complementary tools:

Tool	What it does	Why both
Homepage	A fast, configurable dashboard with live “tiles” for each service	The pretty front door, <i>view</i> status, click through to everything
Portainer	A full web GUI for Docker	The <i>manage</i> half, start/stop/restart containers, read logs, update images, all by clicking

Homepage can *show* container status but can't *control* containers; Portainer *controls* Docker beautifully but isn't a customizable landing page. Together they're the standard beginner homelab combo, and Homepage links to Portainer with a single tile.

i Why Homepage over CasaOS/Cockpit/Homarr

CasaOS wants to own the whole box and install its own Docker stack. It fights the hand-rolled containers you've built so far. **Cockpit** manages the host OS, not Docker, so it's off-target here. **Homarr** is a great click-to-edit alternative if you hate YAML. Homepage wins for a written guide because its config is plain text you can copy-paste, back up, and reproduce exactly.

5.1 Part A: Deploy Homepage and Portainer

5.1.1 Step 1: Deploy Homepage and Portainer

Make the project folder (one stack for the whole control panel):


```
mkdir -p ~/dashboard && cd ~/dashboard
```

Create ~/dashboard/compose.yaml. Note Homepage has **no host port**, Caddy (from Chapter 4) is the front door on port 80 and will reach Homepage by container name over the shared homelab network:

```
services:
  homepage:
    image: ghcr.io/gethomepage/homepage:latest
    container_name: homepage
    volumes:
      - ./config:/app/config
      - /var/run/docker.sock:/var/run/docker.sock:ro
    environment:
      # Hostnames Caddy will forward here. Comma-separated, NO spaces.
      HOMEPAGE_ALLOWED_HOSTS: home.home,homelab,homelab.your-tailnet.ts.net
      # Widget secrets, injected from .env: never hard-coded in the config
      files.
      HOMEPAGE_VAR_PIHOLE_PASSWORD: ${PIHOLE_PASSWORD:?set PIHOLE_PASSWORD
in .env}
      HOMEPAGE_VAR_ABS_TOKEN: ${ABS_TOKEN:-}
    networks:
      - homelab
    restart: unless-stopped

  portainer:
    image: portainer/portainer-ce:lts
    container_name: portainer
    ports:
      - "9443:9443"          # HTTPS management UI (self-signed cert)
    volumes:
      - /var/run/docker.sock:/var/run/docker.sock
      - portainer_data:/data
    networks:
      - homelab
    restart: always

networks:
  homelab:
    external: true          # the shared network created in Chapter 4

volumes:
  portainer_data:
```

Create this project's .env and .gitignore, the same pattern as every other stack:

```
cat > ~/dashboard/.env <<'EOF'
PIHOLE_PASSWORD=same-password-you-set-for-pihole
ABS_TOKEN=your-audiobookshelf-api-token
EOF
chmod 600 ~/dashboard/.env

cat > ~/dashboard/.gitignore <<'EOF'
.env
EOF
```

Get the Audiobookshelf token from its web UI under **Settings** → **Users** → **your user** → **API token**, then bring the stack up:

```
cd ~/dashboard && docker compose up -d
```

💡 Two different “reachability” mechanisms

Homepage uses the **Docker socket** to read container *status* (running/stopped, CPU, RAM), and the **homelab network** to call service *APIs* for widget data. The socket gives you the green “running” dot; the network gives you “1,204 ads blocked today.”

ℹ️ Portainer’s 5-minute clock

The first time you open Portainer (<https://homelab:9443>, accept the self-signed cert warning once), it asks you to create an admin user. Do it **within a few minutes** of starting the container, or Portainer locks itself and you must `docker restart portainer` to reopen the window. The `:lts` image runs natively on a 64-bit Pi.

5.1.2 Step 2: Configure the tiles

❗ Why your dashboard looks empty at first

Homepage does **not** auto-discover anything, out of the box it creates *blank* config files and shows a nearly empty page. Everything you see is whatever you put in `./config/*.yaml`. The blocks below are a deliberately *full* starting point (service tiles, a bookmarks row, and an info bar) so the page looks alive immediately.

Homepage reads YAML from the `./config` folder. Replace the auto-generated blank files with these.

`config/settings.yaml`, look and section layout:

```
title: My Homelab
theme: dark
```

```

color: slate
headerStyle: boxed
layout:
  Media:
    style: row
    columns: 2
  Network:
    style: row
    columns: 2
  Management:
    style: row
    columns: 2
  Links:                # a bookmarks group (defined in bookmarks.yaml
below)
    style: row
    columns: 4

```

`config/docker.yaml`, lets tiles show live container status via the socket:

```

my-docker:
  socket: /var/run/docker.sock

```

`config/services.yaml`, the tiles, now pointing at the pretty URLs from Chapter 4:

```

- Media:
  - Audiobookshelf:
    icon: audiobookshelf.png
    href: https://abs.home                # the pretty URL from
Chapter 4
    description: Audiobooks & language courses
    server: my-docker
    container: audiobookshelf            # -> live status dot +
CPU/RAM
    widget:
      type: audiobookshelf
      url: http://audiobookshelf:80      # container name + internal
port (for widget data)
      key: '{{HOMEPAGE_VAR_ABS_TOKEN}}'  # injected from .env, no
real token in this file

- Network:
  - Pi-hole:
    icon: pi-hole.png
    href: https://pihole.home            # the pretty URL from
Chapter 4
    description: DNS ad-blocking

```

```

        server: my-docker
        container: pihole
        widget:
            type: pihole
            url: http://pihole:80 # container name +
internal port
        version: 6 # REQUIRED for Pi-hole v6
(defaults to 5!)
        key: '{{HOMEPAGE_VAR_PIHOLE_PASSWORD}}' # injected from .env, no
real password in this file

- Management:
  - Portainer:
      icon: portainer.png
      href: https://homelab:9443
      description: Manage Docker containers

```

config/widgets.yaml, the info bar across the top:

```

- resources:
    cpu: true
    memory: true
    disk: /
- search:
    provider: duckduckgo
    target: _blank
- datetime:
    format:
        timeStyle: short
- greeting:
    text_size: xl
    text: Welcome to the homelab

```

config/bookmarks.yaml, a row of quick links (pure links, no Docker logic):

```

- Links:
  - Tailscale admin:
      - abbr: TS
      href: https://login.tailscale.com/admin/machines
  - Router:
      - abbr: RT
      href: http://192.168.1.1
  - Pi-hole docs:
      - abbr: PH
      href: https://docs.pi-hole.net
  - Audiobookshelf docs:

```

```
- abbr: AB
  href: https://www.audiobookshelf.org/docs
```

Apply the config (a plain restart re-reads these files):

```
cd ~/dashboard && docker compose restart homepage
```

Script version for the dashboard stack and config:

```
05-dashboard/scripts/01-create-dashboard-compose.sh
05-dashboard/scripts/02-create-dashboard-env.sh
05-dashboard/scripts/03-create-homepage-config.sh
05-dashboard/scripts/04-start-dashboard.sh
```

The Pi-hole v6 widget gotcha

If the Pi-hole tile shows “API error,” the most common cause is a missing `version: 6` line, the widget defaults to the old v5 API, which no longer exists. The `key` for v6 is your admin password, not a legacy API token.

5.2 Part B: Give it a pretty URL and use it

5.2.1 Step 3: Give the dashboard its pretty URL

This is the Chapter 4 pattern, applied to the dashboard. Add a Caddy route and a Pi-hole record so `https://home.home` (and the alias `https://homelab`) open the dashboard. `tls internal` gives it the same trusted-cert treatment as Chapter 4’s names, no extra certificate work.

1. Add a block to `~/caddy/Caddyfile`:

```
home.home, homelab {
    tls internal
    reverse_proxy homepage:3000
}
```

2. Reload Caddy:

```
cd ~/caddy && docker compose restart caddy
```

3. Add a Pi-hole Local DNS record (Settings → Local DNS Records): `home.home` → your Pi’s Tailscale IP (`100.x.y.z`), exactly like the records you added in Chapter 4.

Script version for the route and DNS record:

```
05-dashboard/scripts/05-add-dashboard-caddy-route.sh
05-dashboard/scripts/06-add-dashboard-dns-record.sh
```

⚠ `HOMEPAGE_ALLOWED_HOSTS` is the #1 “blank page” trap

Homepage refuses to render for any hostname not in its allowed list. Caddy forwards `home.home` and `homelab` to it, so both must appear in `HOMEPAGE_ALLOWED_HOSTS` (they do, in Step 1). If you add another name later, add it there too and `docker compose up -d --force-recreate homepage`. This variable only takes effect on a recreate, not a plain restart.

5.2.2 Step 4: Try it

From your laptop or phone, anywhere on the tailnet:

```
https://home.home      # (or https://homelab, both open the dashboard)
```

You should see your dashboard with live tiles for Audiobookshelf and Pi-hole and a link into Portainer. From the Portainer tile you can start, stop, inspect, and update every container by clicking, including pulling new images.

5.3 Troubleshooting

`https://home.home` shows a blank page. `HOMEPAGE_ALLOWED_HOSTS` doesn't include the host you typed. Add it and `docker compose up -d --force-recreate homepage`.

`https://home.home` times out / “can't connect.” Either the name isn't resolving (check the Pi-hole record from Step 3; try `ping home.home`) or Caddy isn't routing it (confirm the block is in the Caddyfile and you restarted Caddy). `docker logs caddy` shows routing errors.

Widgets show “API error” or stay blank. The widget can't reach the service. Confirm the container is on the `homelab` network (`docker network inspect homelab`) and that you used the container name in the widget url. For Pi-hole, confirm `version: 6` and the correct password.

Portainer won't let me create an admin user. The security timeout elapsed, `docker restart portainer` and reopen `https://homelab:9443` promptly.

5.4 Recap

- **Homepage** (no host port) + **Portainer** (on 9443), both on the `homelab` network, deployed as one `~/dashboard` stack.
- **Tiles** point at the pretty URLs from Chapter 4 and pull live status via the Docker socket and service APIs.

- **Added the dashboard to Caddy**, one Caddy block + one Pi-hole record, so `https://home.home` (and `https://homelab`) open it from anywhere.

Your homelab now has a single, memorable front door. In [Chapter 6](#) we turn to privacy on the *outbound* side: what a VPN does, and how a standard app like Aura compares to Mullvad, then [Chapter 7](#) takes that VPN on the road and explains the Tailscale *exit node*.

Part II

Volume II: VPNs and remote access

6 Chapter 6. Why a VPN, and Mullvad vs. a standard app like Aura

The payoff of this chapter: understand what a VPN does for you, see how a mainstream consumer VPN (I started with Aura) differs from a privacy-focused one (Mullvad), and decide which kind fits this homelab. The plumbing, how a VPN coexists with the tailnet and Pi-hole away from home, is the whole of [Chapter 7](#); this chapter is about choosing the kind of VPN to build around.

Everything so far has been about reaching *into* your homelab. This part is the opposite direction: making your devices' traffic leave through a VPN, so the websites and apps you use can't see your real IP address, and so your ISP or a café's Wi-Fi can't see what you're doing.

Maybe a consumer VPN is already in the mix for changing locations. I started with **Aura VPN**, and it did that job well. I use Aura here as the stand-in for that whole category: NordVPN, ExpressVPN, and the other closed apps built around one big connect button. Mullvad matters in this guide because it can hand over plain WireGuard configs, which gives the homelab something to build with in the next chapter.

6.1 What a VPN does (and what it doesn't)

A VPN routes all your internet traffic through a remote server before it reaches the wider internet. Two practical consequences:

- **Your IP is hidden.** Sites see the VPN server's address, not yours, useful for privacy and for appearing to be in another location.
- **Your local network can't snoop.** On untrusted Wi-Fi, the airport or hotel can see only that you're talking to a VPN, not what you're doing.

What a VPN does *not* do: it isn't anonymity (the VPN provider can see your traffic, so provider trust matters), and it doesn't block ads by itself unless the provider offers DNS-level filtering. Hold onto that last point, because a VPN wanting to own your DNS is where it rubs against Pi-hole, and that collision is the heart of Chapter 7.

6.2 The standard consumer VPN (Aura, and apps like it)

This is the model most people meet first: a self-contained app on each device, one big toggle, maybe a map of countries to pick from. **Aura VPN** is the one I started on, bundled into the Aura identity-protection suite, and it's a fair representative of the whole closed-consumer category (NordVPN, ExpressVPN, and friends behave the same way).

What it does well:

- **It is easy to use.** Install, sign in, tap a country, done. No config files, no command line.
- **It changes your location convincingly** and shields you on sketchy Wi-Fi, which is the common use case.
- **It has split tunneling on its supported platforms**, so you can exempt chosen apps from the tunnel.

Where this category runs into walls for *our* purposes:

- **It's app-only.** There's no client you can run on a headless Raspberry Pi, and no portable config you can lift out and reuse elsewhere. The VPN lives and dies inside the vendor's app.
- **It's a closed box.** You can't see exactly what it does, you trust the vendor's policy, and your VPN identity is tied to the larger account (with Aura, your identity-protection subscription).
- **Like any full-device VPN, it takes over DNS while it's on.** That's normal and even desirable (it prevents DNS leaks), but it's also precisely why it will fight your Pi-hole later.

None of that makes Aura *bad*. It makes it a sealed appliance: useful for “make me look like I’m in another country for an hour,” not built to be reused by a Pi, VPS, or exit-node setup.

6.3 Mullvad: the privacy-first alternative

Mullvad comes at the same job from the opposite philosophy. Where Aura optimizes for a frictionless consumer experience, Mullvad optimizes for privacy and for *control*, and that control is the reason it keeps coming up in this guide.

The points that pushed me toward it:

- **A flat, anonymous price.** A flat **€5/month** (about \$5.40), up to 5 devices, with no tiered “identity suite” wrapped around it. You can even pay by generating an anonymous account number rather than handing over an email.
- **A serious no-logs posture.** Privacy is the product, not an add-on feature, and Mullvad has a public track record (audits, no-logging design) that a bundled consumer VPN generally doesn't lead with.
- **A real kill switch**, so if the tunnel drops, your traffic stops rather than leaking out your real connection.
- **Built-in DNS ad and tracker blocking**, which becomes a useful stand-in on the road when your traffic isn't going through Pi-hole (see Chapter 7).
- **It hands over WireGuard config files.** This is the feature I care about most. Mullvad lets you download standard **WireGuard** configuration files and run them *anywhere*, including on the headless Pi or a cheap Linux VPS, with no Mullvad app at all. Aura gives you nothing like this.

That last point is why Mullvad, not Aura, threads through the rest of the series. A portable WireGuard config can run on another machine, which is what the exit-node setups in Chapter 7 need. The trade is real: less hand-holding, more plumbing to assemble yourself.

i Aura vs. Mullvad

Aura (and consumer VPNs like it) is a sealed app: tap a country and connect, but it stays inside the vendor's client. **Mullvad** is the privacy-first, control-first option: a flat anonymous price, a real kill switch, no-logs by design, and **portable WireGuard config files you can run on any machine**. For casual location-changing, either is fine. For a VPN that needs to mesh with a homelab, Mullvad is the useful shape.

6.4 Which one should you want?

For *this* project, lean Mullvad, for one reason above all the privacy talking points: **it gives you artifacts you can build with**. The portable WireGuard config is what lets a VPN live on your Pi, on a VPS, or inside Tailscale itself. A sealed consumer app like Aura can never do that; it can only protect the one device its app runs on.

Still, keep some perspective:

- If all you ever want is “look like I’m in another country on a phone for an hour,” a consumer app like Aura is fine, and this chapter is enough.
- If you want the VPN to become *part of the homelab*, always-on privacy that still lets you reach your Pi, or a private exit you control, you want Mullvad’s openness, and you want Chapter 7.

6.5 Recap

- A **VPN** hides your IP and shields your traffic from the local network, but it isn’t anonymity and doesn’t block ads on its own.
- The **standard consumer model** (Aura, and apps like it) is sealed and easy: one toggle, one country picker, but app-only, closed, and it can’t run on the Pi or be reused elsewhere.
- **Mullvad** is the privacy-first, control-first alternative: flat anonymous pricing, a real kill switch, no-logs by design, built-in DNS filtering, and **portable WireGuard config files you can run anywhere**.
- **For a homelab, prefer Mullvad**, not mainly for the privacy slogans but because its config files are building blocks; a consumer app is a dead end the moment you want the VPN to mesh with anything else.

Next is the harder part: how that VPN coexists with the tailnet and Pi-hole away from home, what an **exit node** is, and what I run day to day. That’s [Chapter 7](#).

7 Chapter 7. Exit nodes, the Mullvad add-on, and what I run

The payoff of this chapter: the routing rule that governs VPNs on the road, the setup I use (a standalone Mullvad account that I can switch on and off, giving up homelab access while it's on), what a Tailscale **exit node** is, and the paid add-on that *can* give you privacy and your homelab at the same time, if you're willing to trade control and a few more dollars for it.

At home, the pieces share the Pi's network. Away from home, **Tailscale** is the link back to the Pi. This part picks up where [Chapter 6](#) left off: the VPN choice is made; now it has to coexist with the homelab.

7.1 Reaching your homelab from anywhere (the easy part)

This part already works. Because Tailscale is a mesh VPN, any device signed into the tailnet reaches the Pi from anywhere, coffee shop, cellular, hotel, with no extra setup:

- **Audiobookshelf:** <https://abs.home> streams just as it does at home.
- **The dashboard:** <https://home.home> (or <https://homelab>) opens from anywhere.

That is inbound access over an authenticated tunnel. The hard parts below are outbound: where traffic exits and which DNS resolver answers.

7.2 Pi-hole already follows you (set up in Chapter 4)

No extra setup is needed here. When Pi-hole became the tailnet DNS in [Chapter 4](#) (the **Override local DNS** setup), Pi-hole ad-blocking came with it on every tailnet device. On cellular, the phone still resolves through Pi-hole, so ads stay blocked. The same step is also what makes your `.home` names work away from home.

The conflict starts when another VPN is turned on, which is the rest of this chapter.

! This only holds while Tailscale is the active VPN

The whole mechanism rides on Tailscale being the thing controlling your device's DNS. The instant another VPN takes over (next section), this stops applying. That's not a bug, it's the central tension of going mobile, and the rest of this chapter is about navigating it.

7.3 The one rule that governs everything off-network

The road setup follows this rule:

A full-device VPN owns the default route *and* the DNS, and a device runs only one such VPN at a time.

Two consequences fall out of it:

1. **A privacy VPN and Tailscale can't both be the active tunnel.** Phones make this a hard wall: iOS and Android allow exactly one active VPN, so turning on Mullvad (or Aura) drops Tailscale. I ran into the same problem on a laptop, for a different reason: **routing**. A Mullvad WireGuard tunnel grabs the entire default route, which swallows the path back to your tailnet, so the moment Mullvad is up, `https://homelab` and your `100.x` addresses go dark there too. The phone enforces “one VPN” by policy; the laptop enforces it by routing.
2. **Whatever VPN is active forces its own DNS.** Even setting the one-VPN limit aside, the active VPN routes DNS through its own resolvers. So while you're on *any* location/privacy VPN, your traffic is *not* going through Pi-hole, by design, to prevent DNS leaks.

i Why two VPN apps don't give you everything at once

This is why “use my own VPN *and* filter through Pi-hole *and* reach my homelab, all at once, on one device” has no clean answer with two separate apps: Pi-hole and your homelab live only on your tailnet, reaching them needs Tailscale to be the active tunnel, and a privacy VPN that's up takes that role (by policy on a phone, by routing on a laptop). The only real escapes are to put the privacy VPN and your homelab on the *same tunnel*, which is the exit-node idea further down, or to switch between them, which is what I do.

7.4 What I run: keep them separate, switch per activity

After trying to have it all at once and losing (see the sidebar), I settled on the setup that has held up for me: **I keep a standalone Mullvad subscription and Tailscale installed side by side, and I only ever run one at a time.** On both my laptop and my phone.

In practice that means I segregate my activities into two modes and flip between them in a couple of seconds:

- **Tailscale on, Mullvad off (my default).** Homelab reachable, Pi-hole ad-blocking everywhere via the DNS push, `.home` names resolve. This is where I live almost all the time. The cost is that my real IP is showing.
- **Mullvad on, Tailscale off (when privacy matters).** Sketchy Wi-Fi, or I want to change region or hide my IP. While it's on I accept that the homelab and Pi-hole are unavailable for that stretch, Mullvad carries my DNS and its own ad-blocking instead. When I'm done, Mullvad off, Tailscale back on.

My setup is intentionally plain. No exit node, no policy routing, no add-on. The downside is right there in the open: **I can't reach my homelab while the privacy VPN is on**, so I plan around it. For me that beats the alternatives, because I keep full control of my own Mullvad subscription, its config files, its settings, and I'm not paying for anything extra.

i I tried hard to have both on one box, and couldn't make it stick

Before settling on switch-per-activity, I spent hours and hours trying to run Mullvad and Tailscale *simultaneously* on my laptop. The approaches, and why each fell over:

- **Split tunneling / carving out the tailnet.** Mullvad's WireGuard config claims the entire default route (`AllowedIPs = 0.0.0.0/0`), which also swallows the return path to Tailscale's range (`100.64.0.0/10`). The plan was to exclude that range from the Mullvad tunnel so my devices could still reach the homelab.
- **Forcing Tailscale's route priority** so its `100.x` routes would win over Mullvad's catch-all default.

I could occasionally get a fragile version working, but never one that survived a reconnect, a Mullvad server change, or a reboot. It was never reliable enough to trust. The lesson I took: with two separate full-device VPN apps, having both on at once is not a stable setup. If you want both at the same time, change the architecture; that's the exit-node add-on below.

7.5 The thing that *would* let you have both: an exit node

Everything above assumes two separate VPN apps. There's a different architecture that sidesteps the whole "one tunnel at a time" problem, and it's good to know even if, like me, you don't use it yet.

Normal Tailscale is an **overlay network**: it only carries traffic *between* your own devices (laptop homelab phone). Your ordinary internet traffic still leaves through your local connection; Tailscale doesn't touch it.

An **exit node** changes that. You designate one device on your tailnet as "the exit," and your other devices route **all** their internet traffic through it, the full default route, not just tailnet traffic. To the outside world, your traffic now appears to come from the exit node's IP. That's exactly how a traditional full-tunnel VPN behaves, except the "VPN server" is a node on your *own* tailnet, so **you're still on Tailscale the whole time**. Only one exit node is active per device, and you switch it off by selecting "None":

```
tailscale exit-node list          # list available exit nodes
tailscale set --exit-node=<node> # route everything through it
tailscale set --exit-node=       # turn it back off
```

Because an exit node *is* still Tailscale, choosing one doesn't drop your tailnet. Your homelab stays reachable while your traffic exits somewhere else. That architecture escapes the one-rule trap, and it comes in three flavors.

7.6 Flavor 1 (the upgrade path): the Tailscale Mullvad add-on

This is the option that buys you privacy *and* your homelab at the same time, and it's the thing I'd reach for if my needs grow. Tailscale sells access to **Mullvad's worldwide WireGuard servers as a paid add-on purchased through Tailscale**. Mullvad's servers appear as selectable exit nodes inside the tailnet:

```
tailscale exit-node list          # Mullvad cities now appear here
tailscale set --exit-node=<mullvad-node> # exit through Mullvad, stay on
Tailscale
```

On **iOS / macOS / Android**: the **...** menu → **Use exit node** → pick a Mullvad **location**. Because it's still Tailscale, your homelab and `.home` names keep working while your traffic exits through Mullvad. That gives you both pieces: always-on privacy that doesn't cost you the homelab, and it's the only option that works cleanly on a phone without fighting the one-VPN rule.

So why don't I run it? Because of what you give up:

- **Cost:** about **\$5/month** on top of things. If you also keep your own Mullvad subscription (as I do, for the control), you're now at **\$10+/month** for the pair. The add-on alone (~\$5) can *replace* your own sub, but then you lose the next point entirely.
- **Control.** When you buy the add-on, **Tailscale provisions a Mullvad account for you that you have very little say over**. You don't get your own Mullvad login, you can't download WireGuard config files to run on the Pi or a VPS, and you can't tune Mullvad's own settings. You get exit nodes and nothing else. For someone who specifically wants to *own and control* their Mullvad config (which is the whole reason Chapter 6 pointed at Mullvad), that's a real loss.

i Convenience versus control

It comes down to convenience versus control. The add-on gives you *convenience* (one tunnel, always on, homelab intact). Owning your own Mullvad sub gives you *control* (portable configs, full settings, no extra account you don't manage). Right now I value control more, and I'm fine switching apps by hand to get it. If my needs shift toward "always-on privacy that never costs me the homelab," I'll pay the extra and run both, or move fully to the add-on. It's a budget-and-priorities call, not a right-or-wrong one.

When you *do* use a Mullvad exit node, two flags matter on the road:

```
# Keep access to LAN devices (printers, etc.) while using an exit node
tailscale set --exit-node=<mullvad-node> --exit-node-allow-lan-access
```

- **--exit-node-allow-lan-access:** by default, turning on *any* exit node cuts your access to the local network you're physically on. This flag restores it. On mobile it's the **Allow LAN access** toggle when you pick the exit node.
- **DNS leaks:** allowing LAN access can let some DNS queries escape the tunnel. Recent Tailscale clients (1.48.3+) route DNS correctly through Mullvad exit nodes without extra config, so keep your clients updated.

 A real bug to watch for

Some Tailscale client versions around 1.82 had a defect where an **active exit node bypassed the global Pi-hole nameserver**, so your ad-blocking would quietly stop whenever an exit node was on. If you see that symptom, update the Tailscale client first. It's a known issue, not your misconfiguration.

7.7 Flavor 2: your own Pi as an exit node (free, but not private)

You can advertise one of your *own* machines as an exit node. The helper script in this chapter's folder does it on the Pi:

```
sudo tailscale set --advertise-exit-node
# then approve it: admin console -> Machines -> the Pi ->
#   Edit route settings -> enable "Use as exit node"
```

On Linux you also enable IP forwarding (`net.ipv4.ip_forward=1` and the IPv6 equivalent); the script sets both.

- **What it's good for:** routing through your **home** Pi so you appear to be at home, reaching home-only services, or content tied to your home connection, and, importantly, it's the one exit option that keeps your traffic on your own network where **Pi-hole still applies**. This is the sweet spot when what you want on the road is ad-blocking, not IP-hiding.
- **What it's not:** privacy. Traffic exits from **your own home IP**, so there's no hiding. And don't use a *phone* as an exit node, it routes in userspace and is slow; the Pi (kernel routing) is fine.

7.8 Flavor 3 (advanced): a Pi-style exit node chained through Mullvad

Flavor 2 hides nothing, because the box exits from the home IP. You can fix that by *chaining*: run an exit node that itself tunnels out through a standalone Mullvad **WireGuard** config. Your devices reach the node over the tailnet, and the node forwards their traffic out through Mullvad:

you -> Tailscale -> your exit-node box -> Mullvad WireGuard -> internet

This is the only way to use a **standalone** Mullvad subscription *as a Tailscale exit node* without paying for the add-on, so it keeps the control you wanted while still giving you the always-on-plus-homelab shape. The shape of it:

1. Use a machine that is **not your home Pi** if privacy is the goal (a cheap Linux VPS works well). The home Pi only makes sense for "appear at home," where adding Mullvad would defeat the purpose.
2. On that box, install Tailscale, advertise it as an exit node, and enable IP forwarding (same as Flavor 2).

3. Download a **WireGuard config** from your Mullvad account and bring it up (`wg-quick up <conf>`) so the box's default route leaves through Mullvad.
4. The sharp edge: Mullvad's config claims the **entire** default route (`AllowedIPs = 0.0.0.0/0`), which also swallows the return path to your tailnet. You have to keep Tailscale's range (`100.64.0.0/10`) out of the Mullvad tunnel so your devices can still reach the node. That policy-routing fiddliness is exactly the same fight I lost on the laptop, just moved to a box you can leave running, and it's the work the add-on saves you.

Reach for this if you already pay for Mullvad, want to keep your own config, and enjoy the plumbing. For most people the add-on (Flavor 1) buys the same outcome with none of the routing surgery.

7.9 A router-level VPN might be another answer, but I can't test it yet

A friend with 10+ years of homelab experience recently pointed out another possible path: configure the **router itself** to connect to a VPN. In theory, that could move the privacy tunnel out of the phone/laptop app layer and into the network layer, which may avoid some of the conflicts in this chapter.

I can't validate that setup on my current hardware. My router's firmware is too old and does not expose the VPN-client features needed to try it. He also doesn't use Mullvad, so I don't know yet whether this works cleanly with Mullvad's WireGuard configs or whether it introduces a different set of routing or DNS problems.

So treat this as a lead, not a recommendation: if your router supports acting as a WireGuard/OpenVPN client, it may be worth investigating. It is not part of the tested setup in this guide.

7.10 The decision matrix

You can't have all three of {homelab access, Pi-hole filtering, privacy VPN} on one device at one time with two separate apps. But you *can* switch between coherent modes in seconds. Pick per situation:

You're away and you want...	What you turn on	Pi-hole?	Homelab?	Hides IP?
Homelab + ad-blocking (my default)	Tailscale on, no exit node	Yes	Yes	No
Privacy, my way (what I run)	Mullvad on, Tailscale off	No	No	Yes
Appear at home / reach LAN	Tailscale + Pi exit node (Flavor 2)	Yes	Yes	No

You're away and you want...	What you turn on	Pi-hole?	Homelab?	Hides IP?
Privacy <i>and</i> homelab at once	Mullvad add-on exit node (Flavor 1)	No	Yes	Yes

The pattern to read out of that table: with two separate apps (rows 1 and 2) you trade the whole homelab for privacy and switch between them, which is cheap, fully under your control, and what I do. The exit-node rows (3 and 4) are the only ways to keep the homelab *while* your traffic exits elsewhere, and the only one of those that also hides your IP is the paid Mullvad add-on. Two things the Yes/No columns gloss over: in the Mullvad rows your DNS goes to Mullvad, so Pi-hole is bypassed and you lean on Mullvad's own ad-blocking instead; and "appear at home" hides nothing because traffic exits from the home IP.

7.11 A practical everyday routine

For most people the comfortable default is simple, and it's mine:

1. **Leave Tailscale on, no exit node, all the time.** Homelab reachable, Pi-hole ad-blocking everywhere via the DNS push, phone behaves normally. You give up IP-hiding, which most of the time you don't need.
2. **When you specifically want privacy,** switch to standalone Mullvad for that session and accept that the homelab steps aside until you switch back. Cheap, reliable, fully yours.
3. **If "always-on privacy without ever losing the homelab" becomes worth the money and the loss of control,** buy the Mullvad add-on (Flavor 1) and let a Mullvad exit node do both at once, or stand up a chained exit node (Flavor 3) if you'd rather keep your own config.

7.12 Recap: the whole series

- **Reaching your homelab away** is the free win: Tailscale makes `homelab:13378` and `https://home.home` work from anywhere, no extra setup.
- **Pi-hole on the road** rides on the Chapter 4 DNS push, and works only while Tailscale is your active VPN.
- **The governing rule:** a full-device VPN owns the route and DNS, and a device runs one at a time (by policy on a phone, by routing on a laptop), so a separate privacy VPN and your homelab can't both be live with two apps.
- **What I run:** a standalone Mullvad I keep fully under my own control, switched on only when I want privacy, accepting that the homelab is offline while it's on. Boring, reliable, cheap.
- **The escape hatch, when you want both at once:** an **exit node**, most easily the paid **Mullvad add-on**, which trades some control and a few more dollars for always-on privacy that never costs you the homelab.

With that, your homelab is fully mobile: an always-on Raspberry Pi that streams your media, blocks ads on every device, presents a single `https://home.home` control panel, and gives you a deliberate, switchable answer for privacy on the road. All of it private by default, over one Tailscale network, with nothing exposed to the public internet.

That completes the core of the homelab. **Volume III** is the extras: getting a *phone and Linux machines* to talk to each other for everyday remoting, file transfer, and clipboard sharing, starting with [Chapter 8](#).

Part III

Volume III: Extras

8 Chapter 8. Remoting into machines from a phone

The payoff of this chapter: open a terminal on your Pi, or a full graphical desktop on a laptop, straight from a phone, anywhere, by typing the machine's name. No IP addresses, no port forwarding, just the Tailscale name.

By now everything on your tailnet is reachable by name. This chapter puts that to work from a phone. Two tools cover the two common jobs:

Tool	What it gives you	Good for
Termius	An SSH terminal (command line)	The headless Pi, or any server
NoMachine	A full remote desktop (mouse, windows)	A machine that has a GUI (your laptop/desktop)

The Pi has no desktop, so you reach it with **Termius** (SSH). When you want to click around a real desktop from a phone, that's **NoMachine**, pointed at a machine that runs one.

💡 Name your machines first, it makes all of this painless

Everything below addresses machines by their **Tailscale name**, so before anything else, give each device a clear, memorable name. In the Tailscale admin console (<https://login.tailscale.com>) open the **Machines** page, and for each device use the  menu and **Edit machine name** to set something obvious like **homelab**, **laptop**, or **desktop**. With MagicDNS on (Chapter 1), every device on your tailnet can then reach each one by that bare name. Vague auto-generated names like **desktop-3f2a** are the thing that makes remote access annoying; fix that once here.

8.1 Part A. SSH from a phone with Termius

Termius is a free SSH client for iOS and Android. Because the phone runs Tailscale, it resolves your machine names directly, so you connect to **homelab** exactly like you would from your laptop.

8.1.1 Step 1: Confirm the phone is on the tailnet

Open the **Tailscale** app on the phone and make sure it's connected (this was set up in Chapter 1). That's what lets **homelab** resolve. On cellular or any Wi-Fi, the name resolves over the tailnet.

8.1.2 Step 2: Install Termius and add the Pi as a host

Install **Termius** from the App Store or Play Store, then add a new host:

- **Hostname / Address:** the Tailscale name of the machine, e.g. `homelab` (or its full name `homelab.your-tailnet.ts.net`, or the `100.x.y.z` Tailscale IP if a bare name ever doesn't resolve).
- **Username:** your account on that machine (the one you use for `ssh you@homelab`).
- **Port:** 22.

8.1.3 Step 3: Authenticate with an SSH key

Your Pi only accepts **key-based** login (Chapter 1), so Termius needs a key it trusts. Two easy ways:

- **Generate a key in Termius** (Keychain → new key), copy its **public** key, and append it to the Pi's `~/.ssh/authorized_keys` from a machine that's already logged in:

```
echo 'ssh-ed25519 AAAA... (the public key Termius shows you)' >>
~/.ssh/authorized_keys
```

- **Or import the key you already use** from your laptop into Termius's Keychain, since that key is already authorized on the Pi.

Attach the key to the host, tap to connect, and you have a terminal on the Pi from the phone. From here you can run `docker compose`, check logs, restart a service, anything you'd do over SSH at your desk.

8.2 Part B. A full desktop with NoMachine

Sometimes you want a real **graphical desktop**, not a terminal: a browser, a file manager, a GUI app. **NoMachine** streams a machine's desktop to a phone, fast enough to use. It needs a machine that **has** a desktop, so this is for your **Linux laptop or desktop**, not the headless Pi.

8.2.1 Step 4: Install the NoMachine server on the GUI machine

On the machine whose desktop you want to reach (your laptop/desktop), download the NoMachine package for your system from <https://www.nomachine.com/download> and install it. On Debian/Ubuntu that's the `.deb`:

```
sudo dpkg -i nomachine_*.deb      # then, if it complains about deps:
sudo apt --fix-broken install
```

Installing it starts the NoMachine **server**, which listens on port **4000**. No router ports are needed: the phone reaches it over Tailscale.

⚠ Keep NoMachine on the tailnet only

Don't port-forward 4000 from your router. NoMachine should be reachable **only** over your tailnet (and your LAN), exactly like everything else in this series. The Tailscale tunnel is already encrypted; NoMachine adds its own login on top.

8.2.2 Step 5: Connect from the phone by name

Install the **NoMachine** app on the phone, then add a connection:

- **Host:** the machine's Tailscale name, e.g. **desktop** (or its **100.x.y.z**).
- **Port:** 4000.
- **Protocol:** NX.

Log in with that machine's normal user account, and its desktop appears on your phone, controllable by touch. The machine needs to be **on, awake, and logged in to a desktop session** for this to work. When you only need the command line, or you're reaching the Pi, use Termius from Part A instead.

8.3 Recap

- **Name every machine** clearly on the Tailscale Machines page so the rest is just typing a name.
- **Termius** gives you an SSH terminal from a phone, addressed by Tailscale name, with key auth, the right tool for the headless Pi.
- **NoMachine** gives you a full remote desktop from a phone, for machines that run a desktop, over the tailnet with nothing exposed to the internet.

9 Chapter 9. Wiring a phone into Linux machines

The payoff of this chapter: move files between a phone and a Linux laptop in one tap with LocalSend, the app I use on an iPhone, reach your Pi the same way from anywhere over Tailscale, know when KDE Connect’s clipboard sync is worth adding (and when it isn’t), and have the one wired fallback (USB) for when Wi-Fi isn’t an option.

The homelab is now reachable from mobile devices. This part closes the loop in the *other* direction: getting a phone and Linux machines to talk to each other for the everyday stuff, “send me that PDF,” “paste this URL onto my desktop,” “pull these photos off the phone.” On a Mac, AirDrop and Handoff are the obvious defaults; on Linux you assemble the same conveniences from two small apps, plus one wired option you’ll rarely need.

Get one thing straight first: **these are three different tools solving three different problems.** Don’t think “which one wins”, think “which job am I doing.”

9.1 The two Wi-Fi tools (and why I start with LocalSend)

Job	LocalSend	KDE Connect
Send / receive files	Yes	Yes
Share clipboard	No	Yes
Phone battery on desktop	No	Yes
Media remote control	No	Yes
Mirror notifications, SMS, run commands	No	Android only
Dead-simple, one-purpose app	Yes	No

! On an iPhone, I start with LocalSend

KDE Connect’s headline features, notification mirroring, replying to texts from your desktop, running commands on the phone, are **Android-only**: Apple’s sandbox forbids a third-party app from reading other apps’ notifications or sending SMS, so on iOS none of that works no matter what a generic comparison table claims. That leaves the iOS KDE Connect app doing just file transfer, clipboard sync, and battery status, in a clunkier package than LocalSend. For the job almost everyone has on a phone, moving files, **LocalSend is simpler, nicer-looking, more reliable in day-to-day use, and all you need.** Reach for KDE Connect only if you specifically want clipboard sync or battery on your desktop and don’t mind the extra setup; on an iPhone it’s optional, and plenty of people (me included) skip it. One more iOS quirk if you do run it: Apple’s background limits mean the KDE Connect app must be

open or recently used to stay reachable, so it can't sit dormant for days and still answer.

9.1.1 LocalSend: your AirDrop replacement

LocalSend is focused: open it, pick a nearby device, send the file. No pairing, no account, no integration, and it cheerfully moves a 5 GB video that you'd never want to push through a more "integrated" tool.

That's its entire job, and it's the right tool whenever the task is purely "get this file from here to there." The connection finds the other device and the transfer goes through more consistently than KDE Connect's does, and, just as important on an iPhone, **everything the app offers works**, nothing greyed out, Android-only, or quietly broken. That reliability is why I lean on it.

9.1.2 KDE Connect: the optional companion

KDE Connect is the multitool, and on Android it is powerful. On an iPhone its reach is narrower: the one feature worth having is the **shared clipboard** (copy 192.168.1.50 on the phone, paste it on the desktop), with **battery status** on your desktop panel as a minor extra. The notification and SMS superpowers are Android-only, so on iOS that's the whole story, and whether it's worth a second, fiddlier app over just using LocalSend is a real question, answered next.

9.1.3 Do you need KDE Connect on an iPhone?

Probably not. LocalSend already covers the thing most people do constantly, moving files, and it does it with a cleaner app and zero pairing. The only reason to add KDE Connect on an iPhone is if you want **clipboard sync** between phone and desktop often enough to justify a second, clunkier app. If that's you, install both and let KDE Connect handle the clipboard while LocalSend handles files. If it isn't, install LocalSend and stop there, you won't miss anything. (On Android the math flips: KDE Connect's notification and SMS features make it well worth running alongside LocalSend.)

9.2 Install them

9.2.1 On the iPhone

Install **LocalSend** from the App Store, that's the app this chapter starts with. If you also want clipboard sync, grab **KDE Connect** too (the KDE Connect one is published by KDE e.V.); both are free.

9.2.2 On your Linux laptop

LocalSend ships a proper Debian package, so on Ubuntu/Debian the direct install is to grab the .deb from localsend.org/download and hand it to **apt**, which pulls in any dependencies and wires it into your normal updates:

```
# Ubuntu/Debian, run from wherever you saved the file (e.g. ~/Downloads).
# Pick x86-64 for an Intel/AMD laptop, arm-64 for an ARM one.
sudo apt install ./LocalSend-*-linux-x86-64.deb
```

The `./` matters: it tells **apt** to install that local file rather than look up a package by name. On Arch, install **localsend-bin** from the AUR. (A distro-agnostic AppImage is also on the [download page](#) if you'd rather not install anything.)

KDE Connect, from your distro's repos:

```
# Arch
sudo pacman -S kdeconnect

# Ubuntu/Debian
sudo apt install kdeconnect
```

i Running KDE Connect on Hyprland (no Plasma, no GNOME)

KDE Connect does **not** require the full KDE Plasma desktop, but it does need its background daemon running and a way to interact with it. On a bare window manager like Hyprland:

- The daemon **kdeconnectd** is started on demand over D-Bus; launching the GUI (**kdeconnect-app**) or running **kdeconnect-cli -l** brings it up.
- For an at-a-glance tray/battery indicator, add a **Waybar** “custom” module that calls **kdeconnect-cli**, or just keep **kdeconnect-app** a keybind away.
- Skip **GSConnect**, it's a GNOME Shell extension and won't help on a bare Wayland setup like this. Plain **kdeconnect** is the right choice here.

9.2.3 Open the firewall

Both tools need their ports reachable on the LAN. If you run **ufw**:

```
sudo ufw allow 53317/udp      # LocalSend discovery + transfer
sudo ufw allow 53317/tcp
sudo ufw allow 1714:1764/udp  # KDE Connect
sudo ufw allow 1714:1764/tcp
```

If a transfer “sees” the device but hangs at 0%, a closed firewall port is the first thing to check.

9.2.4 Pair

- **LocalSend:** no pairing, open it on both devices, on the same Wi-Fi, and they appear in each other's list. Send.
- **KDE Connect:** open the app on both, tap the laptop in the phone's device list, and **accept the pairing request** on the desktop (`kdeconnect-app` or the CLI will prompt). Pairing is a one-time trust handshake.

9.3 Making it work *anywhere*: over Tailscale

Here's the homelab tie-in. Out of the box, both tools find devices using **LAN multi-cast/broadcast discovery**, great at home, useless the moment your phone is on cellular, because broadcast packets don't traverse a VPN. But your tailnet gives every device a stable `100.x.y.z` address, and both apps let you **add a device by IP manually**, which sidesteps discovery entirely:

- **LocalSend:** Settings → **add a favorite / manual device** → enter the target's Tailscale IP (e.g. your laptop at `100.x.y.z`) with port 53317. Now you can send to your laptop from the airport.
- **KDE Connect:** the iOS app has an “**Add device by IP**” / refresh option, enter the laptop's Tailscale IP. Pair once, and clipboard + file transfer work across the tailnet.



Why this is the natural payoff of everything before it

This is the same trick that powered the whole series: discovery and “same network” assumptions break the instant you leave home, and Tailscale's stable addresses are the fix. Just as you reached Audiobookshelf and `https://home.home` by tailnet IP, you reach your *laptop* by tailnet IP here. The phone Linux conveniences stop being a home-only luxury.

9.4 Phone the headless Pi

Keep one thing in mind here: LocalSend and KDE Connect are **desktop** tools. Your Pi from Chapters 1–5 runs a **headless** server OS (Ubuntu Server 26.04 LTS), no graphical desktop, so neither app is the natural fit there. For moving a file between a phone and the Pi, you already have better, GUI-free paths:

- **Drop it into a service you already run.** Files you want *on* the Pi (more audiobooks, say) can go straight into `~/audiobookshelf/media/Audiobooks`, `scp` them over the tailnet:

```
# from your laptop, over Tailscale, no LAN required
scp ~/Downloads/lesson.mp3
you@homelab:~/audiobookshelf/media/Audiobooks/
```

- **From the phone directly**, an iOS file-manager app that speaks SSH/SFTP (several exist) can connect to **homelab** over Tailscale and drop files onto the Pi without any desktop app in between.

In short: **laptop phone** is LocalSend/KDE Connect territory; **phone/laptop the headless Pi** stays SSH/scp territory, which you already have for free over Tailscale.

9.5 The wired fallback: USB access with libimobiledevice

Everything above is Wi-Fi. USB is a **completely separate technology stack**, and most people rarely need it. Install it when one of these comes up, not before:

- No Wi-Fi available and you need photos off the phone *now*.
- You’re moving 100 GB of video and want the fastest possible transfer.
- The phone’s networking is broken and you need emergency access.

When that day comes:

```
# Arch
sudo pacman -S libimobiledevice ifuse usbmuxd

# Ubuntu/Debian
sudo apt install libimobiledevice6 libimobiledevice-utils ifuse usbmuxd
```

Then plug in, trust the computer on the phone, and mount:

```
idevicepair pair                # tap "Trust" on the iPhone when prompted
mkdir -p ~/iphone
ifuse ~/iphone                  # mounts the phone's media (camera roll)
# ... copy your files ...
fusermount -u ~/iphone          # unmount when done
```

i What USB gives you (and what it doesn’t)

usbmuxd is the daemon that talks to the phone over the USB cable; **libimobiledevice** is the library implementing Apple’s protocols; **ifuse** mounts the result as a folder. Because of iOS’s sandbox, **ifuse** exposes the **photos/ camera-roll area** (and, with extra flags, individual apps’ document folders), *not* the whole iOS filesystem. So USB is excellent for bulk photo/video offload, but it isn’t a “browse my entire iPhone like a USB stick” experience. No jailbreak, no full filesystem, that’s an Apple limitation, not a missing package.

9.6 Where to start

On an iPhone, begin with one app:

1. **LocalSend**, the AirDrop replacement that handles the file job everyone has, and whose features work reliably.

Add **KDE Connect** only if you find yourself wanting clipboard sync between phone and desktop; on Android, add it from the start, where its notification and SMS features earn their keep. Either way, set up the Tailscale manual-IP pairing once you want these to work away from home. Then, *if* months later you find yourself wanting to pull a whole camera roll off over a cable, install **libimobiledevice** + **ifuse** at that point, when you'll know exactly why you need them, instead of carrying tools you never touch.

For most people, Wi-Fi transfer (LocalSend / KDE Connect, extended over Tailscale) handles the overwhelming majority of phone Linux moments. USB is a backup and a power-user tool, not a daily requirement.

9.7 Recap: and the series, complete

- **LocalSend** = start here: dead-simple file send with a clean app whose features work reliably, install it on the example iPhone and the laptop and you're set for the common case.
- **KDE Connect** = optional companion: handy for clipboard sync and battery, but on an iPhone it's limited (notifications/SMS are Android-only) and clunkier, so add it only if you want clipboard sync, it earns its place mainly on Android.
- **Tailscale extends both off-LAN**: add devices by their 100.x.y.z IP so transfer and clipboard work from anywhere, exactly as the rest of your homelab does.
- **The headless Pi** stays on SSH/scp over Tailscale, the right tool for a machine with no desktop.
- **USB (libimobiledevice + ifuse)** is the wired fallback for bulk photo offload or no-Wi-Fi emergencies, install it only when needed.

That completes *Beginner Homelab on a Raspberry Pi*. You started by building a foundation (a hardened, always-on Pi on a private Tailscale network) and then gave it job after job: streaming your media, blocking ads on every device, clean `http://*.home` URLs behind a reverse proxy, a single `https://home.home` control panel, a deliberate answer for VPN privacy on the road, and now two-way file sharing with a phone. Every piece runs on hardware under the reader's control, over one private network, with nothing exposed to the public internet.

A Appendix A. The Docker Compose files, line by line

The main chapters keep the `compose.yaml` files terse on purpose. You don't need to understand every line to get a working homelab. This appendix is the opposite: it explains **every directive** in every Compose file the series creates, so when you come back in six months you can read your own stack and know exactly what each line does (and what's safe to change).

By the end of the series your Pi runs five containers across four Compose projects:

Project folder	File	Containers
~/audiobookshelf	compose.yaml	audiobookshelf
~/pihole	compose.yaml	pihole
~/caddy	compose.yaml	caddy
~/dashboard	compose.yaml	homepage, portainer

A.1 Compose concepts that show up everywhere

Before the per-file walkthroughs, here are the directives you'll see repeatedly. Read this once and the individual files become obvious.

- **services:**, the top-level map. Each key under it (e.g. `pihole:`) is one container Compose will manage.
- **image:**, which prebuilt image to pull and run, in `repository:tag` form. `pihole/pihole:latest` means "the latest tag of the `pihole/pihole` image." `:latest` floats; pinning a version (e.g. `:2024.07.0`) is more reproducible but you update by hand. This guide uses `:latest` for simplicity and updates with `docker compose pull`.
- **container_name:**, the fixed, human-friendly name (e.g. `pihole`). Without it, Compose auto-generates names like `pihole-pihole-1`. We set it explicitly because **the reverse proxy and the inter-container networking address each service by this exact name**, `reverse_proxy pihole:80` only works if the container is literally named `pihole`.
- **ports:**, publishes a **host** port to a **container** port, written `"HOST:CONTAINER"`. `"8081:80"` means "traffic arriving on the Pi's port 8081 is forwarded to port 80 inside the container." A bare `"80:80"` binds **all** the host's network interfaces (LAN *and* tailnet); prefixing an address, `"192.168.1.50:80:80"`, binds only that one interface. Ports are only needed for traffic entering from *outside* Docker, containers on the same Docker network reach each other directly without any **ports:** entry.
- **environment:**, variables set inside the container. This is how most images are configured. Values can be literals or `${VAR}` references resolved from the project's `.env` file (see below).

- **volumes:**, persistent storage. Two forms appear in this series:
 - **Bind mount**, `./etc-pihole:/etc/pihole` maps a *folder on the Pi* to a path inside the container. You can see and back up the files directly. The `./` is relative to the folder the `compose.yaml` lives in. A `:ro` suffix (`...:/etc/caddy/Caddyfile:ro`) mounts it **read-only**.
 - **Named volume**, `portainer_data:/data` stores data in a Docker-managed volume (declared in the top-level `volumes:` block). Use it when you don't need to touch the files yourself, only persist them across container recreations.
- **networks:**, which Docker networks the container joins. Containers on the same network resolve each other **by container name** as a hostname. This is the backbone of the reverse proxy: Caddy and the services share the `homelab` network, so `pihole`, `audiobookshelf`, and `homepage` are resolvable names.
- **restart:**, the restart policy. `unless-stopped` restarts the container on crash and on boot, *unless* you deliberately stopped it. `always` is similar but restarts even one you stopped, after a Docker daemon restart. Both are correct for an always-on Pi; the difference rarely matters.
- **Top-level networks:** / **volumes:**, declare resources the services reference. `external: true` means “this network already exists; don't create or destroy it, just attach to it”, used for the shared `homelab` network that outlives any single project.

i `${VAR}` interpolation, `.env`, and the `?:` / `:-` forms

Compose substitutes `${VAR}` from a file named `.env` in the same folder *before* starting the container. This is the mechanism that keeps secrets out of the file you commit:

- `${PIHOLE_PASSWORD}`, plain substitution; empty if unset.
- `${PIHOLE_PASSWORD:?set PIHOLE_PASSWORD in .env}`, **required:** Compose aborts with that error message if the variable is missing or empty. Used for secrets you must not launch without.
- `${ABS_TOKEN:-}`, **default if unset:** the part after `:-` is the fallback (here, empty string), so a missing optional token doesn't error.

The real values live in `.env`, which is listed in `.gitignore` and never leaves the Pi. The `compose.yaml` holds only references, so it's safe to publish.

A.2 ~/pihole/compose.yaml: Pi-hole

```
services:
  pihole:
    container_name: pihole
    image: pihole/pihole:latest

    ports:
      - "53:53/tcp"
      - "53:53/udp"
```

```

    - "8081:80/tcp"

environment:
  TZ: "${TZ}"
  FTLCONF_webserver_api_password: "${PIHOLE_PASSWORD:?set
PIHOLE_PASSWORD in .env}"
  FTLCONF_dns_upstreams: "1.1.1.1;1.0.0.1"
  FTLCONF_dns_listeningMode: "ALL"

volumes:
  - ./etc-pihole:/etc/pihole

networks:
  - homelab

restart: unless-stopped

networks:
  homelab:
    external: true

```

- **ports: 53:53/tcp and 53:53/udp**, DNS uses **both** transports (UDP for almost everything, TCP for large responses), so both are published. They’re bound to all interfaces so the Pi answers DNS for your LAN *and* over the tailnet. (In Chapter 3, before Caddy existed, these were bound to the Pi’s LAN IP only, `${PIHOLE_IP}:53:53`; Chapter 4 widens them.)
- **ports: 8081:80/tcp**, Pi-hole’s web UI listens on port 80 *inside* the container; we publish it on the host as **8081** so the host’s port 80 stays free for Caddy.
- **TZ**, the timezone, so Pi-hole’s logs and graphs use your local time.
- **FTLCONF_***, Pi-hole v6’s configuration system. Each variable maps to a key in `/etc/pihole/pihole.toml` by the rule `FTLCONF_<section>_<key> → [section] key`. Anything set this way is **re-applied on every container start** and shows as read-only in the web UI, exactly what you want for file-defined, reproducible config:
 - `FTLCONF_webserver_api_password`, the admin/API password (the v6 successor to the old `WEBPASSWORD`).
 - `FTLCONF_dns_upstreams`, the real resolvers Pi-hole forwards *non-blocked* queries to; `1.1.1.1;1.0.0.1` is Cloudflare’s primary + secondary.
 - `FTLCONF_dns_listeningMode: "ALL"`, answer queries on any interface. Required for a Dockerized Pi-hole, because LAN clients’ queries arrive *through* the Docker gateway and the stricter “local-only” mode would reject them.
- **volumes: ./etc-pihole:/etc/pihole**, persists Pi-hole’s entire state (config, blocklists, query database, your local DNS records) in a folder next to the compose file, so recreating the container loses nothing.
- **networks: homelab + top-level external: true**, joins the shared homelab network *declaratively*. This is deliberate: attaching it by hand with `docker network connect` would be wiped by the next `--force-recreate` (see the warning in [Chapter 4](#)). In the file, it always re-attaches.

A.3 ~/caddy/compose.yaml + Caddyfile: the reverse proxy

```
services:
  caddy:
    container_name: caddy
    image: caddy:latest
    ports:
      - "80:80"
      - "443:443"
      - "443:443/udp"
    volumes:
      - ./Caddyfile:/etc/caddy/Caddyfile:ro
      - caddy_data:/data
      - caddy_config:/config
    networks:
      - homelab
    restart: unless-stopped

networks:
  homelab:
    external: true

volumes:
  caddy_data:
  caddy_config:
```

- **ports:** 80:80, 443:443, 443:443/udp, Caddy is the **single front door**, so no other container may publish these. Port **80** catches plain-HTTP requests and redirects them to HTTPS; port **443** serves HTTPS (the /udp line enables HTTP/3). Every *.home request lands here first.
- **volumes:**
 - ./Caddyfile:/etc/caddy/Caddyfile:ro, mounts your routing table into the container, **read-only** (Caddy never needs to write its own config).
 - caddy_data / caddy_config, named volumes for Caddy's internal state. Crucially, caddy_data holds the **internal certificate authority** that tls internal creates (/data/caddy/pki/...) and the certs it signs, so your trusted CA survives container recreations. Kept as named volumes because you never edit them by hand.
- **networks:** homelab, lets Caddy reach each backend by container name. This is the whole trick: Caddy connects to pihole:80, audiobookshelf:80, homepage:3000 *over this network*, completely ignoring the host port mappings. That's why moving Pi-hole's host UI port to 8081 doesn't affect Caddy.

The companion **Caddyfile** is the routing table:

```

pihole.home {
    tls internal
    redir / /admin 302
    reverse_proxy pihole:80
}

abs.home {
    tls internal
    reverse_proxy audiobookshelf:80
}

home.home, homelab {
    tls internal
    reverse_proxy homepage:3000
}

```

- **{ ... } site block**, one per hostname (or comma-separated list of hostnames). The address is a **bare name** (no `http://`/`https://`); Caddy matches the incoming request's `Host` header to the right block.
- **tls internal**, the key directive. There's no public certificate authority for a private `.home` name, so Caddy spins up its **own** local CA, signs a cert for each name, and serves HTTPS. You trust that CA once per device (Chapter 4 Step 7) and the URLs get a real padlock. Caddy also auto-redirects `http://` to `https://` for these names.
- **reverse_proxy <name>:<port>**, forwards the request to that container over the `homelab` network. The port is the service's *internal* port, not its published host port.
- **redir / /admin 302**, Pi-hole's UI lives under `/admin`; this sends a bare `https://pihole.home/` to `https://pihole.home/admin` so you land in the right place. The 302 is a temporary-redirect status code.
- **home.home, homelab**, two names, one backend: both open the dashboard.

A.4 ~/dashboard/compose.yaml: Homepage + Portainer

```

services:
  homepage:
    image: ghcr.io/gethomepage/homepage:latest
    container_name: homepage
    volumes:
      - ./config:/app/config
      - /var/run/docker.sock:/var/run/docker.sock:ro
    environment:
      HOMEPAGE_ALLOWED_HOSTS: home.home,homelab,homelab.your-tailnet.ts.net
      HOMEPAGE_VAR_PIHOLE_PASSWORD: ${PIHOLE_PASSWORD:?set PIHOLE_PASSWORD
in .env}
      HOMEPAGE_VAR_ABS_TOKEN: ${ABS_TOKEN:-}

```

```

    networks:
      - homelab
    restart: unless-stopped

portainer:
  image: portainer/portainer-ce:latest
  container_name: portainer
  ports:
    - "9443:9443"
  volumes:
    - /var/run/docker.sock:/var/run/docker.sock
    - portainer_data:/data
  networks:
    - homelab
  restart: always

networks:
  homelab:
    external: true

volumes:
  portainer_data:

```

This is the only project with **two** services. Note Homepage has **no ports:** , it's reached only through Caddy (homepage:3000 over homelab), so it never touches a host port.

Homepage:

- **image:** `ghcr.io/gethomepage/homepage:latest`, the full registry path; this one lives on GitHub's Container Registry (`ghcr.io`) rather than Docker Hub.
- **volumes:**
 - `./config:/app/config`, your dashboard's YAML config (tiles, widgets, bookmarks) on the Pi, editable directly.
 - `/var/run/docker.sock:/var/run/docker.sock:ro`, the **Docker socket, read-only**. This lets Homepage read container status (the live "running" dot, CPU/RAM) without being able to control Docker.
- **Homepage_ALLOWED_HOSTS**, a required allow-list of hostnames Homepage will render for. It must contain *every* name Caddy forwards here (`home.home`, `homelab`, and the Pi's full tailnet name). A request for a name not on the list gets a blank page, the single most common Homepage gotcha. **This variable only takes effect on a container recreate** (up `-d --force-recreate`), not a plain `restart`.
- **Homepage_VAR_***, Homepage exposes any `Homepage_VAR_<NAME>` variable to its config files as `{{Homepage_VAR_<NAME>}}`. This is how widget secrets (the Pi-hole password, the Audio-bookshelf API token) reach the config **without being written into the config files**. They come from `.env` at runtime.

Portainer:

- **image:** `portainer/portainer-ce:lts`, Community Edition, Long-Term Support tag. Runs natively on the 64-bit Pi.
- **ports:** `9443:9443`, Portainer's own HTTPS UI (self-signed certificate). It keeps a host port because you reach it directly at `https://homelab:9443`, not through Caddy.
- **volumes:**
 - `/var/run/docker.sock:/var/run/docker.sock`, the Docker socket, **read-write this time** (no `:ro`). Portainer's whole job is to *control* Docker, start, stop, recreate, pull, so it needs full socket access.
 - `portainer_data:/data`, a named volume for Portainer's own database (users, settings).
- **restart:** `always`, Portainer is your break-glass management console; you want it back even after a manual stop + daemon restart.

 The Docker socket is root-equivalent

Mounting `/var/run/docker.sock` into a container gives that container control over Docker, which on Linux is effectively root on the host. That's why Homepage gets it `:ro` (it only needs to *read* status) and only Portainer gets it read-write (it needs to *manage*). Never expose a container that mounts the socket to the public internet; on your tailnet-only homelab it's fine.

A.5 ~/audiobookshelf/compose.yaml: Audiobookshelf

The media server set up in [Chapter 2](#). It's the first stack you create, then [Chapter 4](#) adds the `networks` block so Caddy can reach it as `audiobookshelf:80`:

```
services:
  audiobookshelf:
    container_name: audiobookshelf
    image: ghcr.io/advplyr/audiobookshelf:latest

    ports:
      - "13378:80"

    volumes:
      - ./media/Audiobooks:/audiobooks
      - ./config:/config
      - ./metadata:/metadata

    networks:
      - homelab

    restart: unless-stopped

networks:
  homelab:
    external: true
```

The whole stack is self-contained in `~/audiobookshelf/`: the `compose.yaml`, the app's `config/` and `metadata/`, and the content under `media/`. All the volume paths are relative to that folder.

- **ports:** `13378:80`, Audiobookshelf listens on port 80 *inside* the container; we publish it on the host as **13378** (host 80 is reserved for Caddy). You still reach the app directly at `homelab:13378`, the iPhone app uses that address, while the browser uses `abs.home` through Caddy.
- **volumes:**
 - `./media/Audiobooks:/audiobooks`, your audiobook library. A future Podcasts library is just another folder, `./media/Podcasts:/podcasts`. (Audiobookshelf is for spoken audio; music gets its own server.)
 - `./config:/config` and `./metadata:/metadata`, Audiobookshelf's database (users, listening positions, settings) and cached cover art, kept next to the compose file so recreating the container loses nothing.
- **networks:** `homelab + top-level external: true`, joins the shared homelab network *declaratively*, exactly like Pi-hole. Added in Chapter 4; before that the stack runs on its own default network. Declaring it here (rather than a hand `docker network connect`) is what makes the attachment survive every recreate, see the warning under Pi-hole above.

B Appendix B. What should just work, and how to verify it

This is the at-a-glance reference for a finished build: every URL, what it's for, and the one-line check that proves it's healthy. Use it as a smoke test after a reboot, after an update, or any time something feels off, work top to bottom and the first failing row tells you which layer broke.

Throughout, substitute your own values where this guide uses placeholders: the Pi's Tailscale IP for `100.x.y.z` (run `tailscale ip -4` on the Pi) and your tailnet suffix for `your-tailnet.ts.net`.

B.1 The one-time cloud setting everything depends on

Almost everything below is configured *on the Pi* and works the moment the containers are up. The **single exception** is a setting in the Tailscale admin console that can't be done from the Pi, and until it's on, the `.home` names and tailnet-wide ad-blocking won't work on your *other* devices:

! Make Pi-hole your tailnet's DNS (admin console, one time)

In the Tailscale admin console (<https://login.tailscale.com>) → **DNS** page:

1. **Add nameserver** → **Custom**, enter the Pi's **Tailscale** IP (`100.x.y.z`), leave *Restrict to domain* **off**, and *Use with exit node* **off**. Save.
2. Turn on **“Override local DNS.”**

This is the [Chapter 4](#) Step 5 step. Verify it took with `tailscale dns status` on any device, under *Resolvers* you should see your Pi's `100.x.y.z`, not “no resolvers configured.” If it still says no resolvers, the admin-console setting hasn't applied yet.

B.2 The service map

Service	URL (everywhere on your tailnet)	What it does
Dashboard (Homepage)	<code>https://home.home</code> (or <code>https://homelab</code>)	Landing page; live tiles + links
Audiobookshelf	<code>https://abs.home</code>	Stream your audiobooks / Assimil
Pi-hole	<code>https://pihole.home</code>	DNS ad-blocking admin
Portainer	<code>https://homelab:9443</code>	Docker management GUI

Service	URL (everywhere on your tailnet)	What it does
---------	----------------------------------	--------------

All four work from any device signed into your tailnet, home Wi-Fi or cellular, with no IPs and no port numbers (except Portainer, which keeps its own :9443).

B.3 Verify it, layer by layer

Run these in order. Each isolates one layer, so the first failure pinpoints the problem.

1. The containers are all up (on the Pi):

```
docker ps --format "table {{.Names}}\t{{.Status}}"
```

You should see five **Up** containers: `caddy`, `homepage`, `pihole`, `portainer`, `audiobookshelf`.

2. They share the homelab network (on the Pi):

```
docker network inspect homelab --format "{{range .Containers}}{{.Name}}
{{end}}"
```

All five names should appear. If `pihole` or `audiobookshelf` is missing, the reverse proxy can't reach it, re-read the network callout in [Chapter 4](#).

3. DNS resolves the pretty names (on the Pi):

```
dig +short pihole.home @127.0.0.1      # expect your Pi's Tailscale IP
(100.x.y.z)
dig +short doubleclick.net @127.0.0.1 # expect 0.0.0.0 (ad-blocking works)
```

4. Caddy routes each name to the right service (on the Pi):

```
for name in pihole.home abs.home home.home; do
    printf "%-12s -> " "$name"
    curl -s -o /dev/null -w "HTTP %{http_code}\n" \
        --resolve "$name:80:$(tailscale ip -4 | head -1)" "http://$name/"
done
```

Expect `pihole.home` -> HTTP 302 (it redirects to `/admin`) and `abs.home` / `home.home` -> HTTP 200. This tests the full chain, DNS name → Caddy on the Tailscale IP → the correct backend container, exactly as a remote device sees it.

5. The names work from another device (laptop/phone on the tailnet):

```
ping home.home          # answers from the Pi's 100.x.y.z
```

Then open `https://home.home`, `https://abs.home`, and `https://pihole.home` in a browser. If ping fails here but step 4 passed on the Pi, the missing piece is the **Tailscale DNS override** (the cloud setting at the top of this appendix).

6. The dashboard widgets have live data:

Open `https://home.home`. The Pi-hole tile should show today’s blocked-query count and the Audiobookshelf tile your library, proof the widgets authenticated through the **homelab** network using the secrets in `~/dashboard/.env`. If a tile says “API error,” see the widget troubleshooting in [Chapter 5](#) (for Pi-hole, the usual culprit is a missing `version: 6`).

B.4 What works away from home (and what deliberately doesn’t)

Once the cloud setting above is on, this is what happens the moment you leave the house, the full reasoning is in [Chapter 7](#):

Mode (away)	Homelab URLs?	Pi-hole ad-block?	Hides your IP?
Tailscale on, no exit node (everyday default)	Yes	Yes	No
Tailscale + your Pi as exit node	Yes	Yes	No (home IP)
Tailscale + Mullvad exit node	Yes	No (Mullvad DNS)	Yes
Aura (or any standalone VPN) on	No	No	Yes

The one rule behind the table: a phone runs **one** VPN at a time, and whatever VPN is active owns DNS. So homelab access, Pi-hole filtering, and a privacy VPN can’t all be on at once. You switch between coherent modes. The everyday default (top row) is the one you live in.

B.5 Reboot behavior

Nothing extra is required after a power cycle. Every container uses a `restart: policy` (`unless-stopped` or `always`), so Docker brings them all back on boot; the **homelab** network is persistent; and the `systemd-resolved` stub stays disabled (the drop-in from [Chapter 3](#) survives reboots), so Pi-hole reclaims port 53 cleanly. Reboot the Pi and re-run the verification steps above, every row should pass without you touching anything.